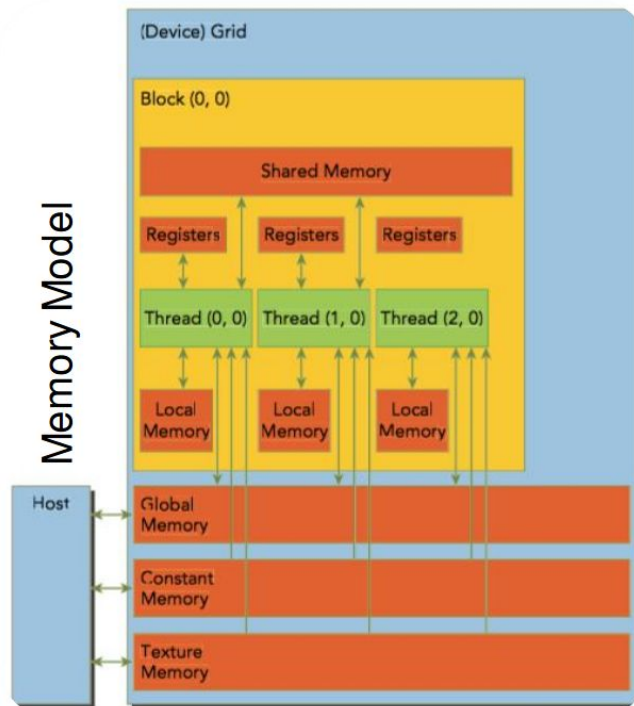


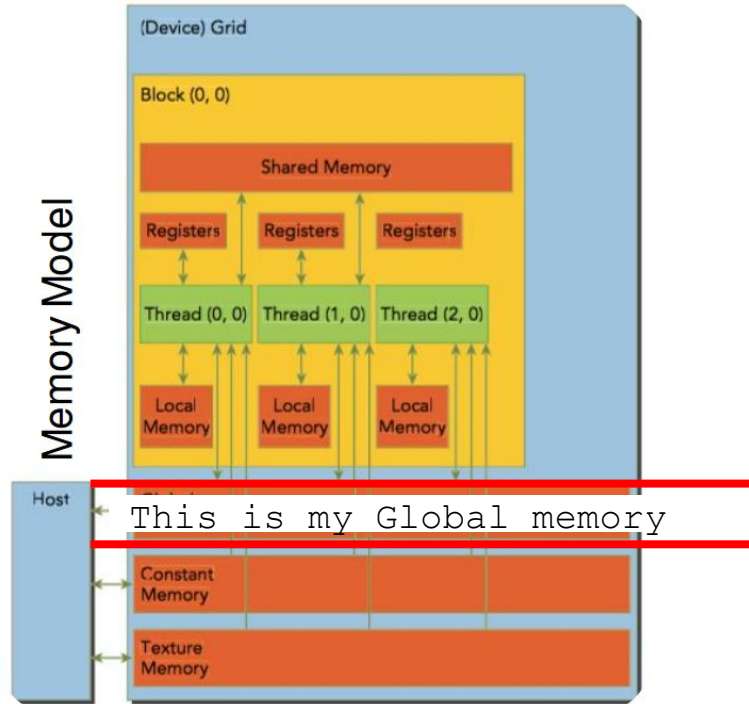
Global memory and Coalescing

A click through example
By Nicolas Bohm Agostini

Memory Model




Memory Model

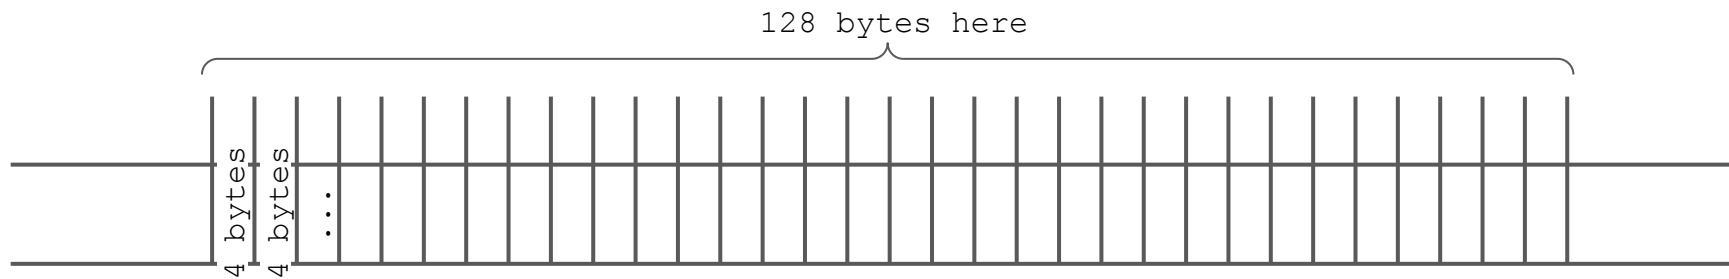


This is my Global memory

This is my Global memory




These are groups of
4 bytes



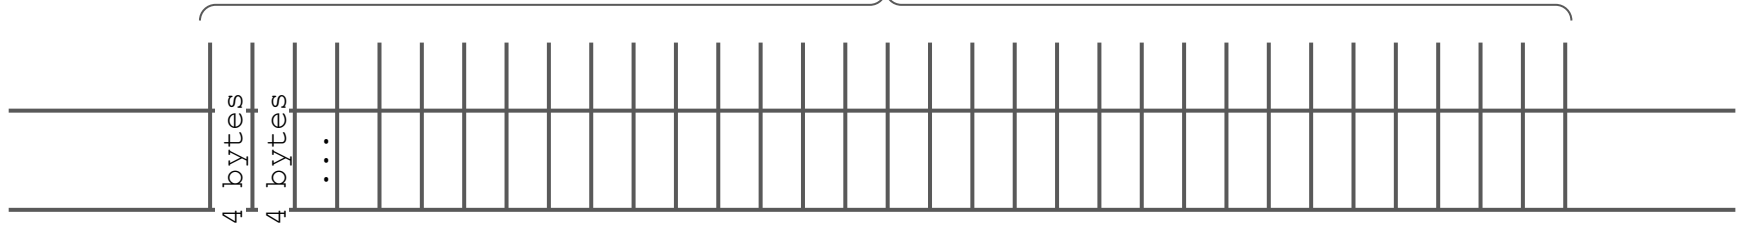
These are groups of
4 bytes

128 bytes here. 32 Groups of 4 bytes



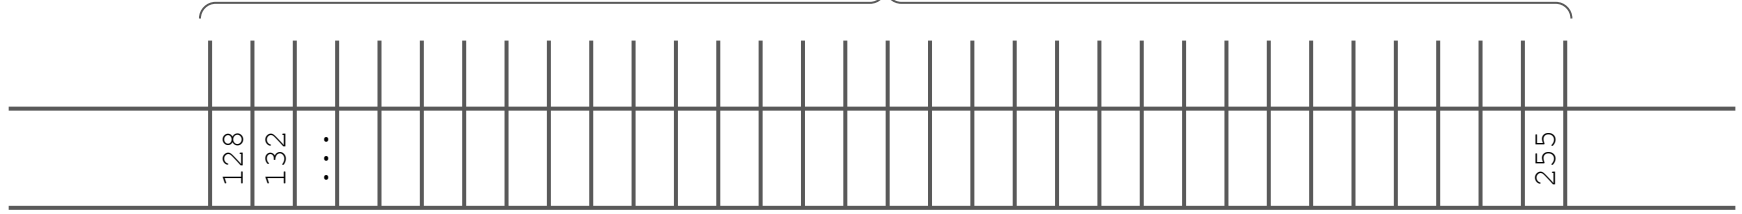
These are groups of
4 bytes

128 bytes here. 32 Groups of 4 bytes



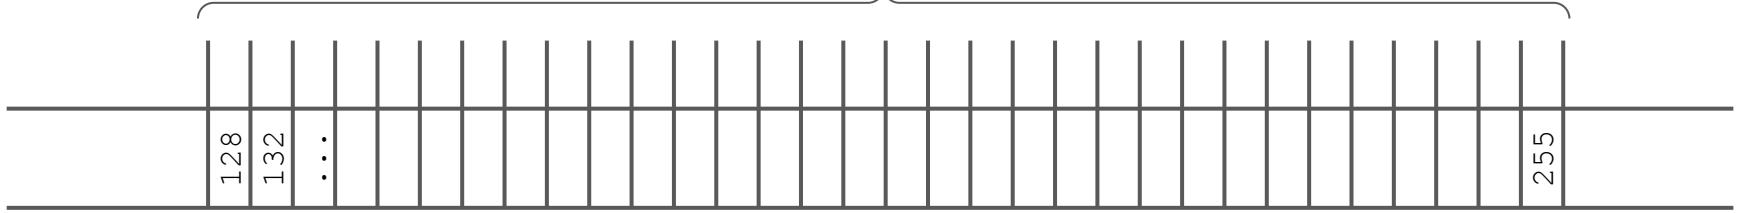
These are groups of
4 bytes

128 bytes here. 32 Groups of 4 bytes



This chunk of data got allocated starting on byte 128 of my
Global Memory

128 bytes here. 32 Groups of 4 bytes



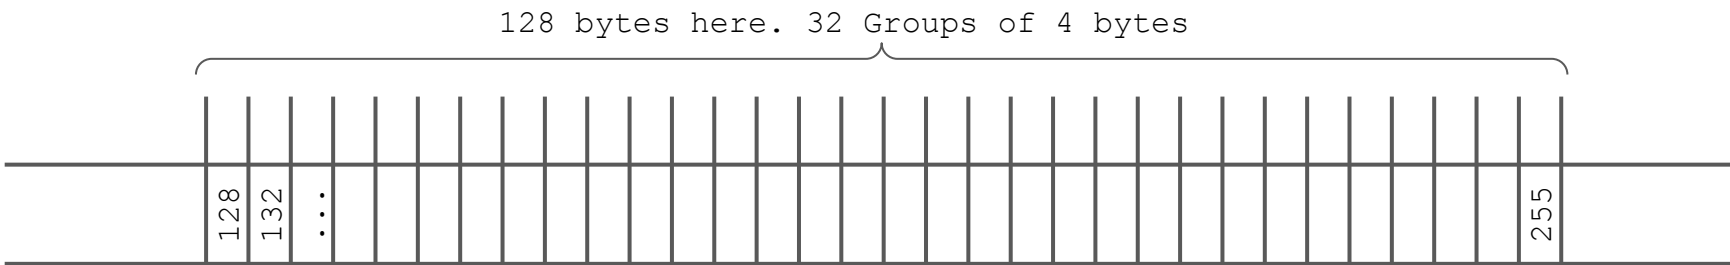
```
__device__ void add1(int* a, int vector_size) {
    id = blockIdx.x*blockDim + threadIdx.x;

    if (id<vector_size)
        a[id]=a[id]+1;
}

int main() {
    // do setup

    add1<<<4,256>>>(dev_a,vector_size);


    // do cleanup
}
```



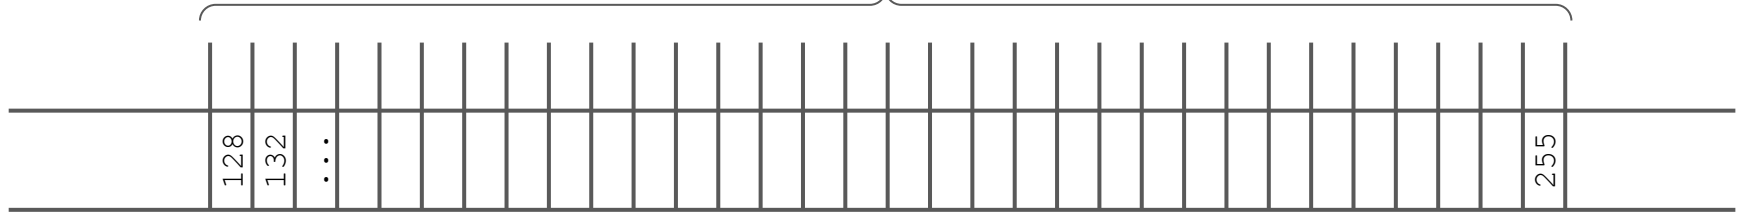
```
int main() {
    // do setup

    add1<<<4,256>>>(dev_a,vector_size);

    // do cleanup
}
```



128 bytes here. 32 Groups of 4 bytes




```
int main() {  
    // do setup  
  
    add1<<<4,256>>>(dev_a,vector_size);  
  
    // do cleanup  
}
```

...

block

...

128 bytes here. 32 Groups of 4 bytes



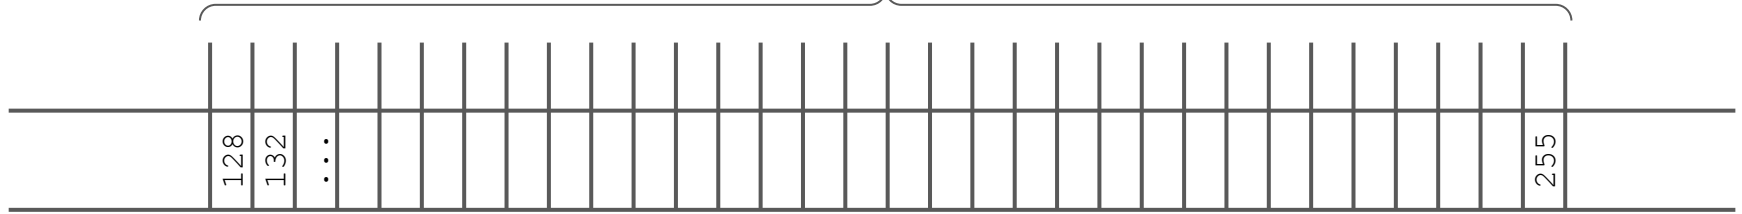
```
int main() {  
    // do setup  
  
    add1<<<4,256>>>(dev_a,vector_size);  
  
    // do cleanup  
}
```

...

block with 256/32 warps

...

128 bytes here. 32 Groups of 4 bytes




```
int main() {  
    // do setup  
  
    add1<<<4,256>>>(dev_a,vector_size);  
  
    // do cleanup  
}
```

...

block with 8 warps

...

128 bytes here. 32 Groups of 4 bytes



```
int main() {  
    // do setup  
  
    add1<<<4,256>>>(dev_a,vector_size);  
  
    // do cleanup  
}
```

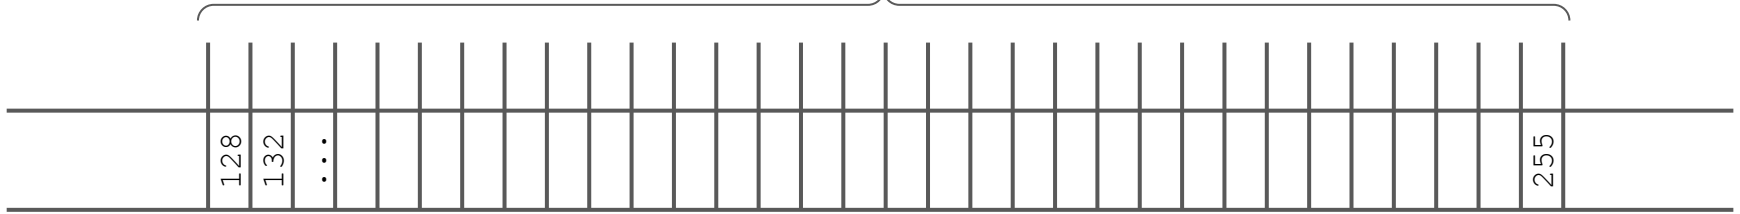
1 WARP: 32 threads

...

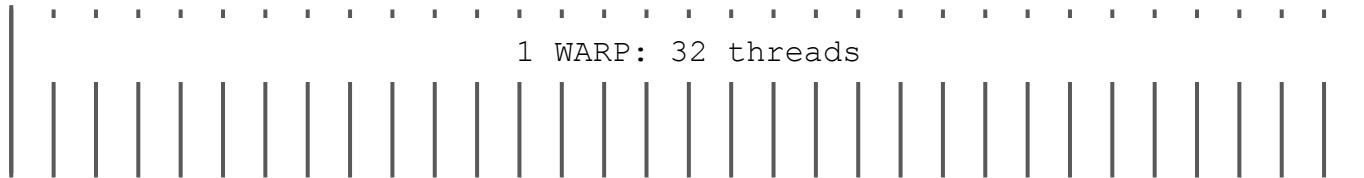
block with 8 warps

...

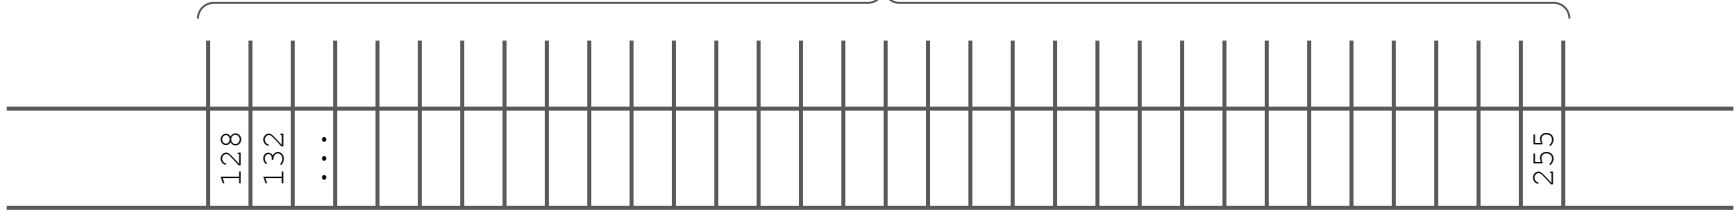
128 bytes here. 32 Groups of 4 bytes



```
int main() {  
    // do setup  
  
    add1<<<4,256>>>(dev_a,vector_size);  
  
    // do cleanup  
}
```

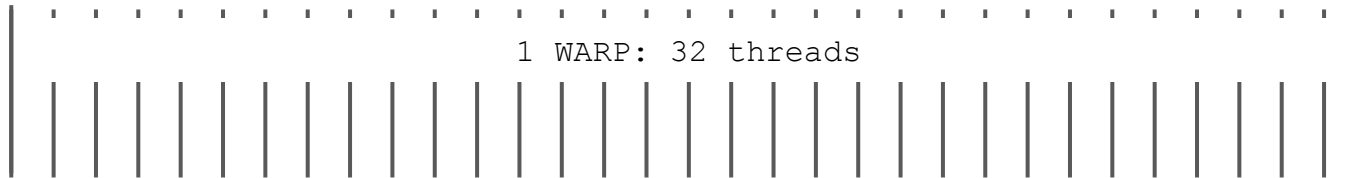


128 bytes here. 32 Groups of 4 bytes




```
__device__ void add1(int* a, int vector_size) {  
    id = blockIdx.x*blockDim + threadIdx.x;  
  
    if (id<vector_size)  
        a[id]=a[id]+1;  
}
```

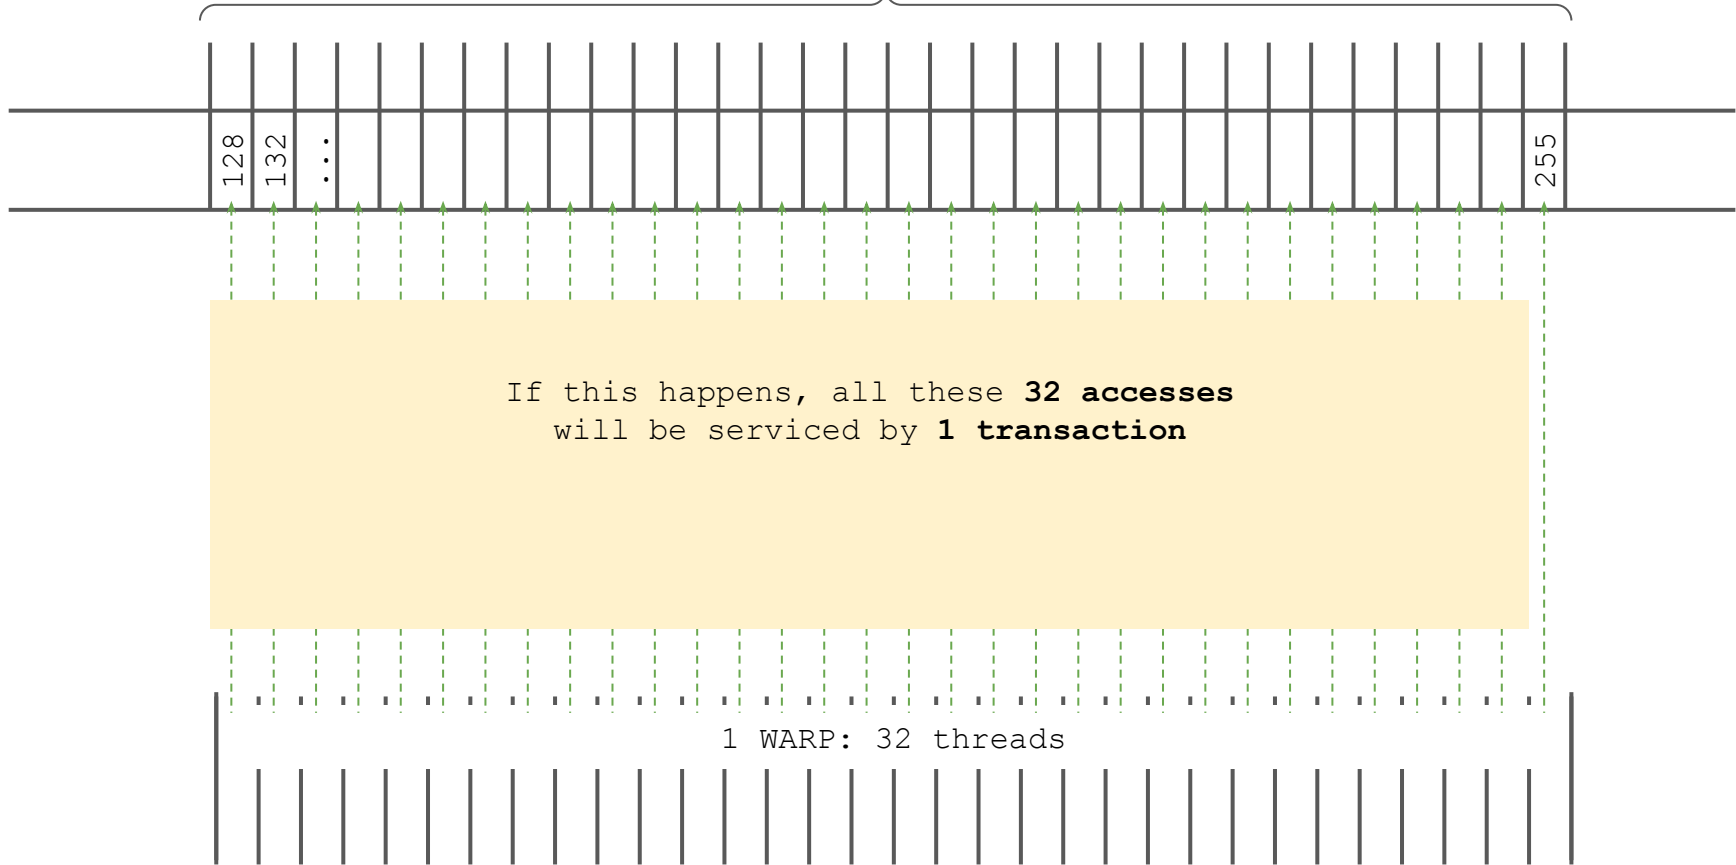
Each thread reads one element
of array **a**



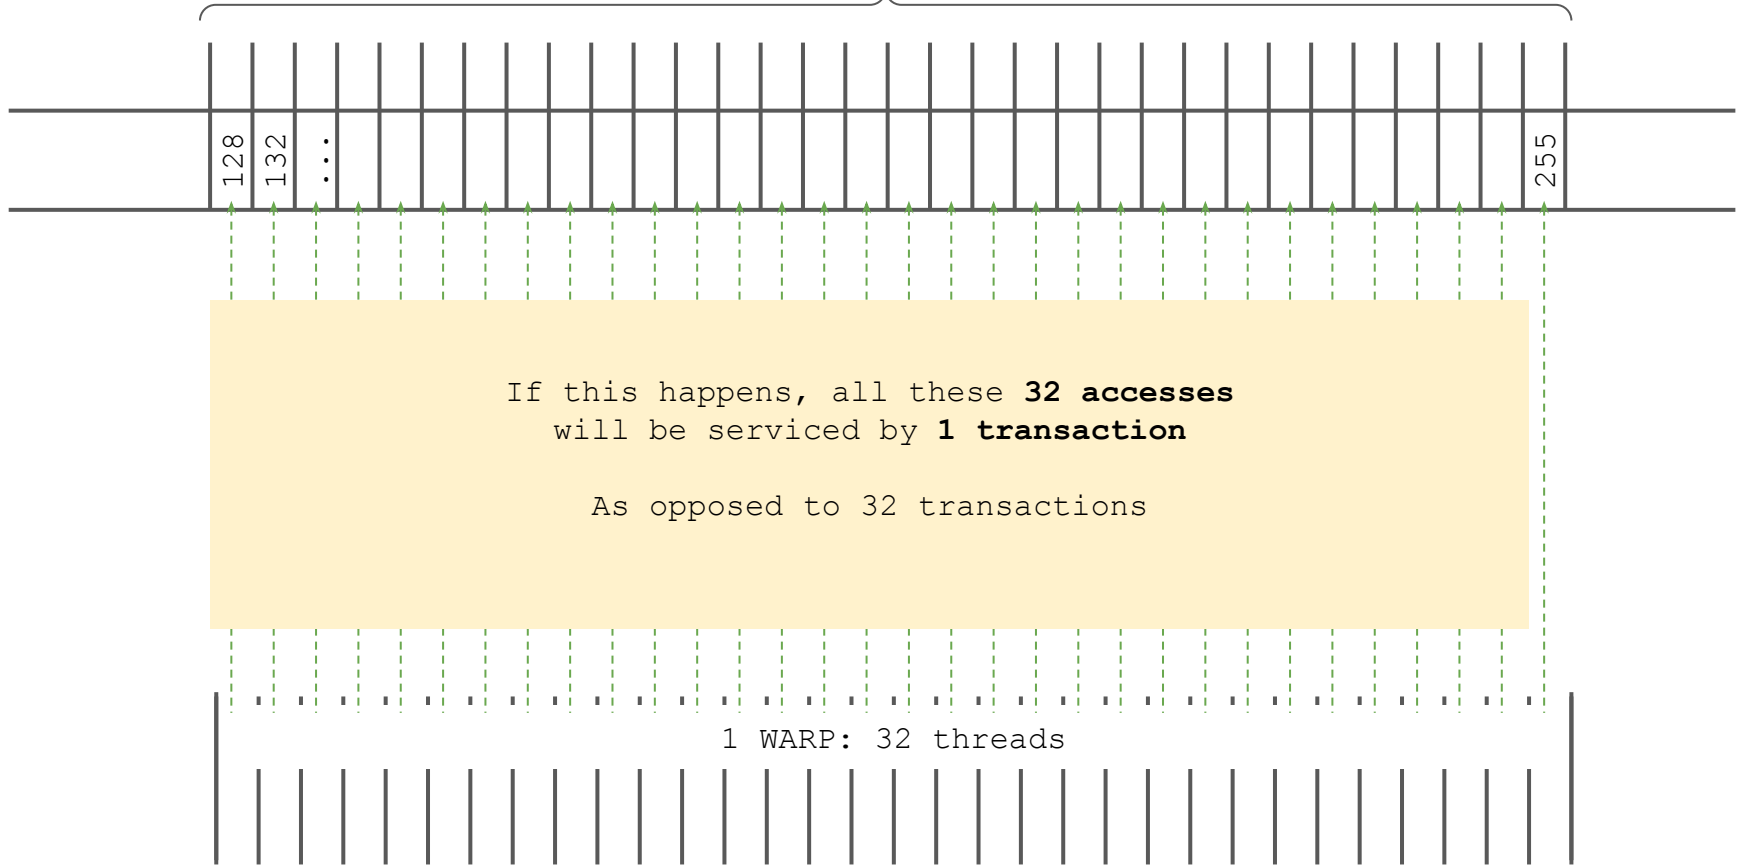
128 bytes here. 32 Groups of 4 bytes



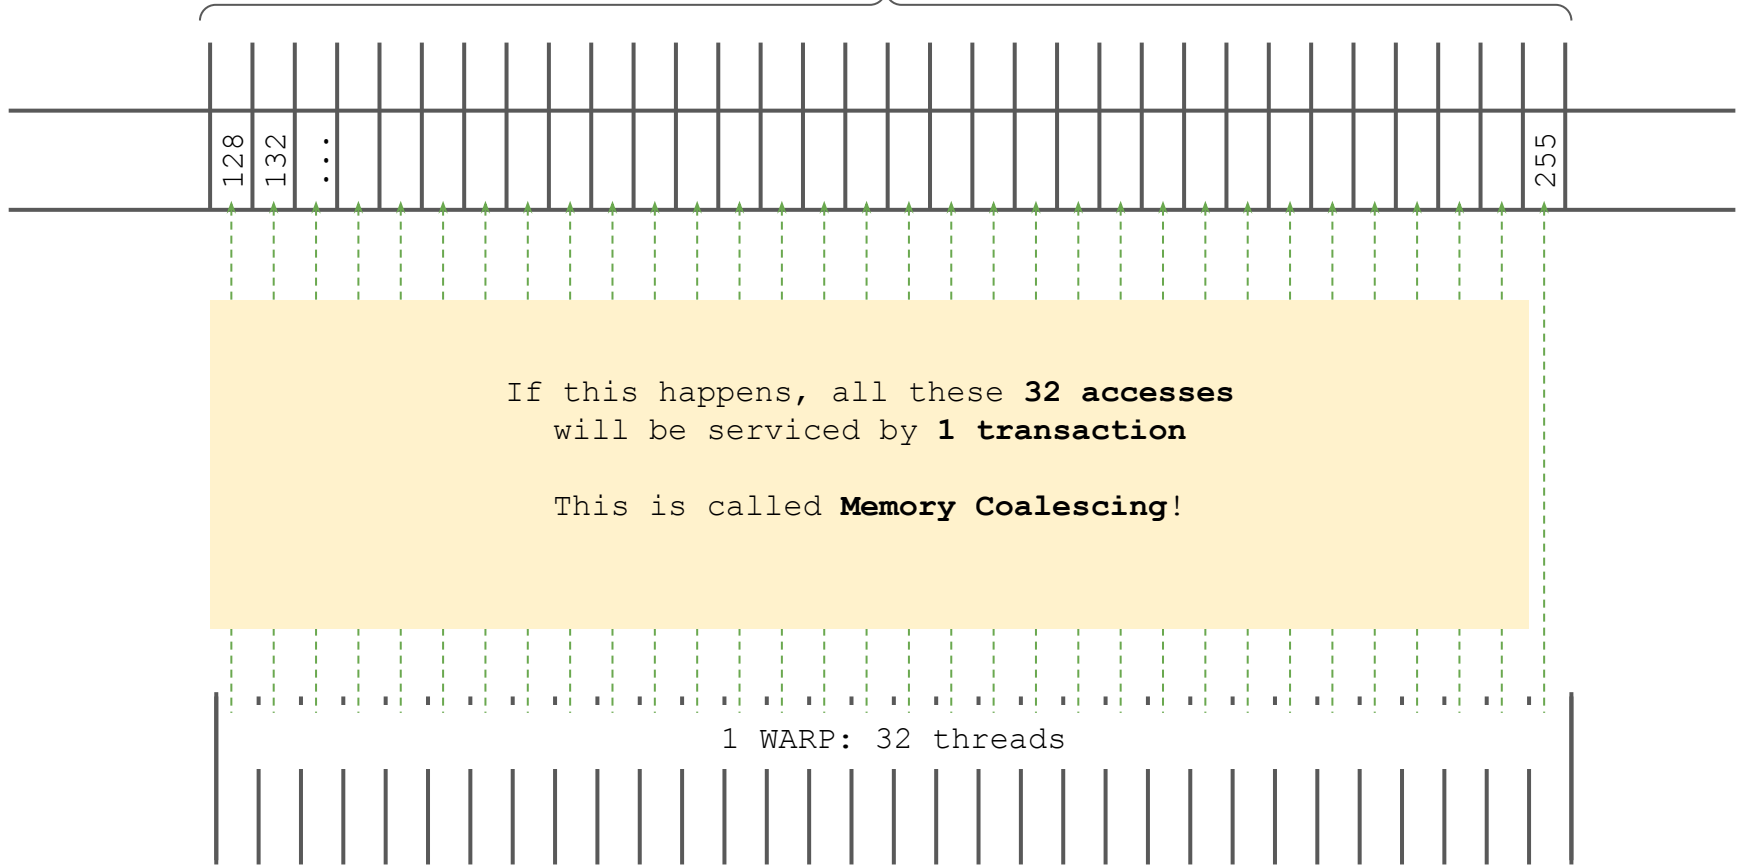
128 bytes here. 32 Groups of 4 bytes



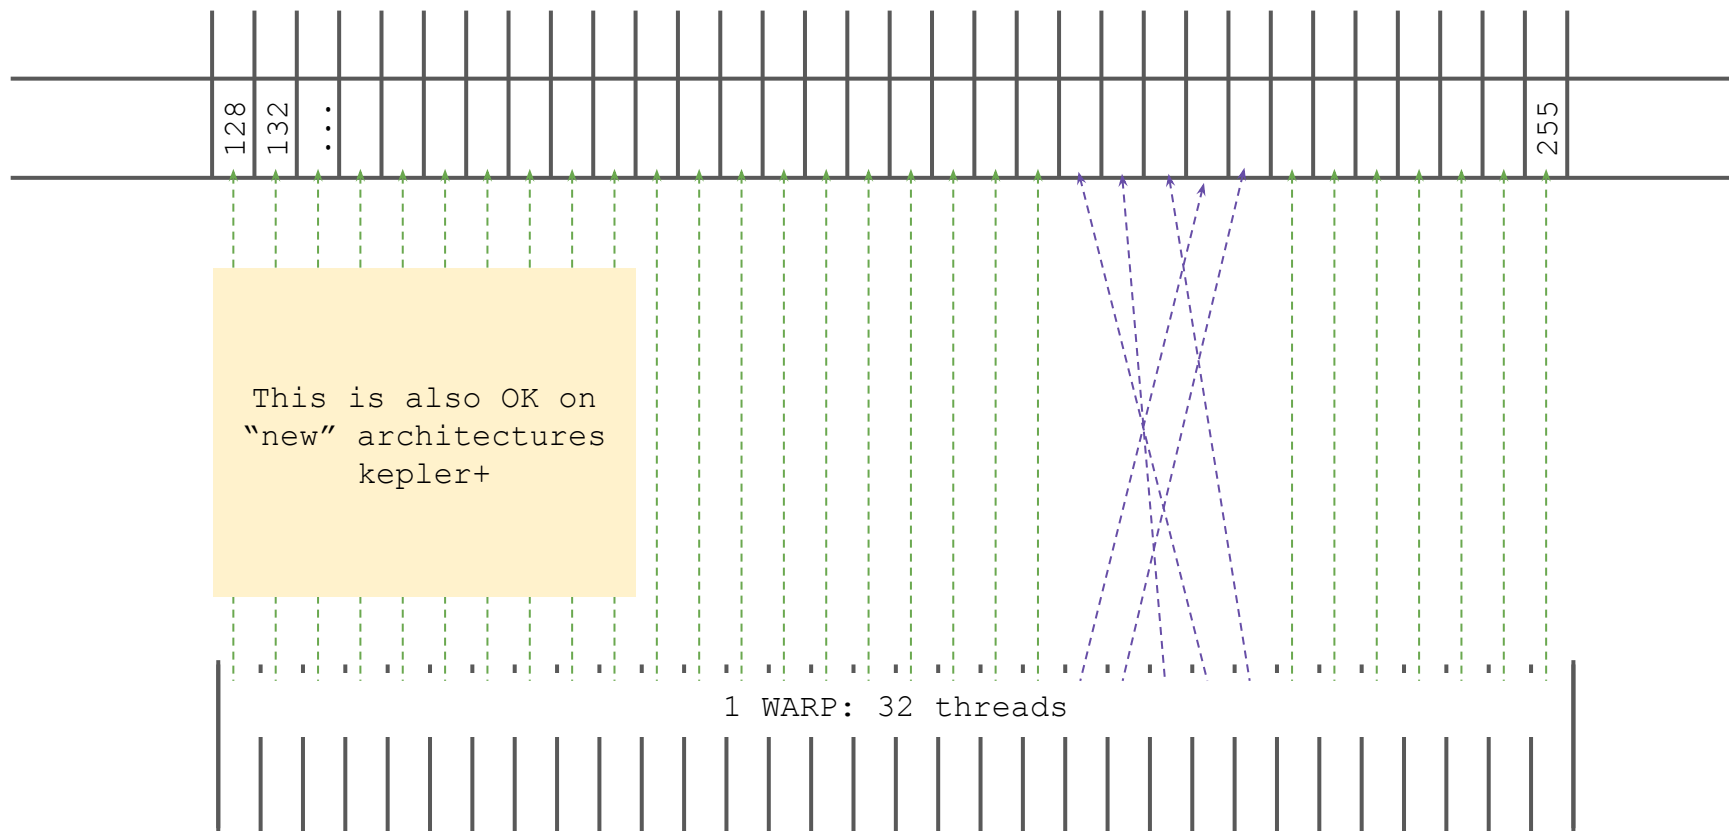
128 bytes here. 32 Groups of 4 bytes



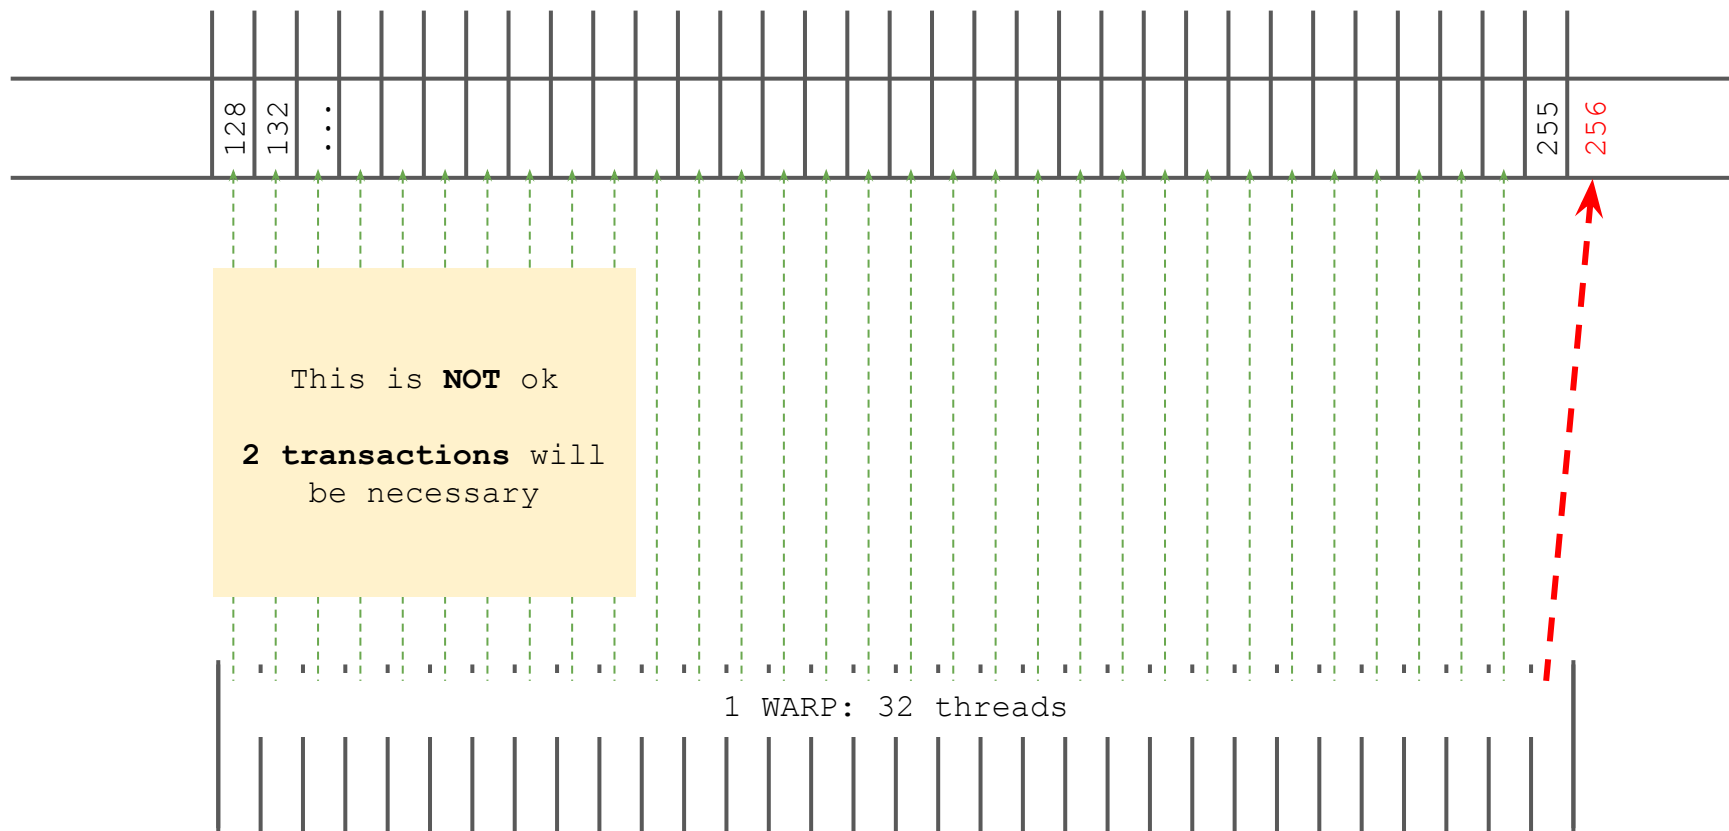
128 bytes here. 32 Groups of 4 bytes



Another kernel... that has some non-contiguous access based on tid



Yet Another kernel... that has some non-contiguous access based on tid



In your lab today!

```
Typedef struct myCoordinates {  
    int x;  
    int y;  
    int z;  
} Coordinates;
```

```
int main () {  
    // do stuff  
  
    coordinates * a = malloc(sizeof(Coordinates)*vector_size);  
  
    // do more stuff  
}
```

In your lab today!

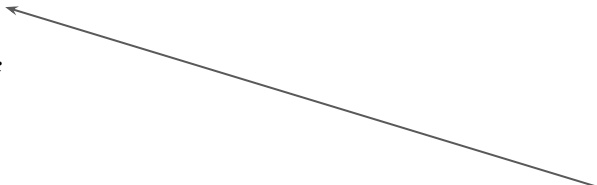
```
typedef struct myCoordinates {  
    int x;  
    int y;  
    int z;  
} Coordinates;
```

```
int main () {  
    // do stuff
```

```
    coordinates * a = malloc(sizeof(Coordinates)*vector_size);
```

```
    // do more stuff
```

```
}
```



Size of 3 integers:
12 bytes

In your lab today!

```
Typedef struct myCoordinates {  
    int x;  
    int y;  
    int z;  
} Coordinates;
```

```
int main () {  
    // do stuff  
  
    coordinates * a = malloc(sizeof(Coordinates)*vector_size);  
  
    // do more stuff  
}
```

In your lab today!

```
typedef struct myCoordinates {  
    int x;  
    int y;  
    int z;  
} Coordinates;
```

In memory 1 Coordinates
object would look like this

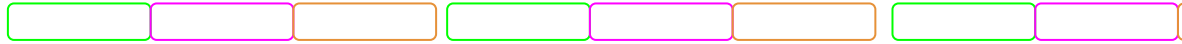


```
int main () {  
    // do stuff  
  
    coordinates * a = malloc(sizeof(Coordinates)*vector_size);  
  
    // do more stuff  
}
```

In your lab today!

```
typedef struct myCoordinates {  
    int x;  
    int y;  
    int z;  
} Coordinates;
```

In memory Several
Coordinates object would
look like this

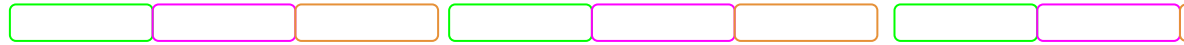


```
int main () {  
    // do stuff  
  
    coordinates * a = malloc(sizeof(Coordinates)*vector_size);  
  
    // do more stuff  
}
```

In your lab today!

```
typedef struct myCoordinates {  
    int x;  
    int y;  
    int z;  
} Coordinates;
```

So, accessing only **x (the greens)** would be a strided access



```
int main () {  
    // do stuff  
  
    coordinates * a = malloc(sizeof(Coordinates)*vector_size);  
  
    // do more stuff  
}
```

In your lab today!

```
Typedef struct myCoordinates {  
    int x;  
    int y;  
    int z;  
} Coordinates;
```

What we want is:



```
int main () {  
    // do stuff  
  
    coordinates * a = malloc(sizeof(Coordinates)*vector_size);  
  
    // do more stuff  
}
```

Good luck! ヽ_(ツ)_/