



GPU Programming

(in CUDA)

Summer 2020

Designed by Julian Gutierrez, Presented by Nicolas Agostini

Session 11





- Improving Host to Device Memory Transfers
 - Pinned Memory
 - Unified Memory
- Concurrency
- Dynamic Parallelism

Improving Host to Device Memory Transfers

Memory Transfers



- We want to minimize the amount of time it takes to transfer data between host and device.
- Device to host memory bandwidth is much lower than device to device bandwidth

Memory Transfers



- We want to minimize the amount of time it takes to transfer data between host and device.
 - How can we do this?
 - Reduce the amount of data we have to transfer! (if possible)
 - Achieve highest memory bandwidth between host and device. (Possible when using page-locked (or “pinned”) memory, we are going to talk about it in this session)
 - Batching many small transfers into one larger transfer performs much better because it eliminates most of the per-transfer overhead
 - Data transfers between the host and device can sometimes be overlapped with kernel execution and other data transfers (we will talk about this in this session as well)
-


- Remember
 - When deciding whether to run on the GPU or on the CPU, we should consider the kernel execution time and the time it takes to copy data back and forth.
 - We need to consider the cost of moving data across the PCI-e bus.
-

Pinned Memory (Page-Locked Data Transfers)




- Host (CPU) data allocations are pageable by default.
 - The GPU cannot access data directly from pageable host memory.
 - Due to this, the Cuda driver must:
 - Allocate a temporary page-locked block, or “pinned”
 - Copy host data to the pinned block
 - Transfer from pinned to device
 - Delete pinned block

Pageable Data Transfer



Pinned Data Transfer




Pinned Memory (Page-Locked Data Transfers)




- `cudaMallocHost()`
 - Prevents OS from paging host memory
 - Allows PCI-e DMA to run at full speed
- `cudaFreeHost()`
- Allocations can fail so you need to check errors.

Pageable Data Transfer



`malloc`

Pinned Data Transfer



`cudaMallocHost`

```
// allocate and initialize
h_aPageable = (float*)malloc(bytes); // host pageable
h_bPageable = (float*)malloc(bytes); // host pageable
checkCuda( cudaMallocHost((void**)&h_aPinned, bytes) ); //
host pinned
checkCuda( cudaMallocHost((void**)&h_bPinned, bytes) ); //
host pinned
checkCuda( cudaMalloc((void**)&d_a, bytes) ); // device
```

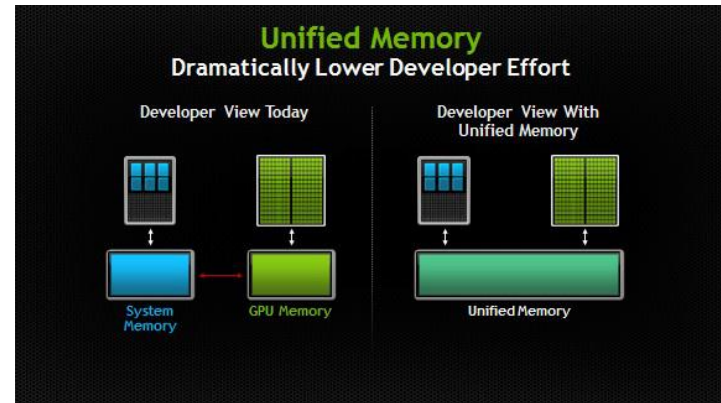

Pinned Memory (Page-Locked Data Transfers)

- You should not over-allocate pinned memory. Doing so can reduce overall system performance because it reduces the amount of physical memory available to the operating system and other programs.
 - Transfers between the host and device are the slowest link of data movement involved in GPU computing, so you should take care to minimize transfers.
-

Unified Memory



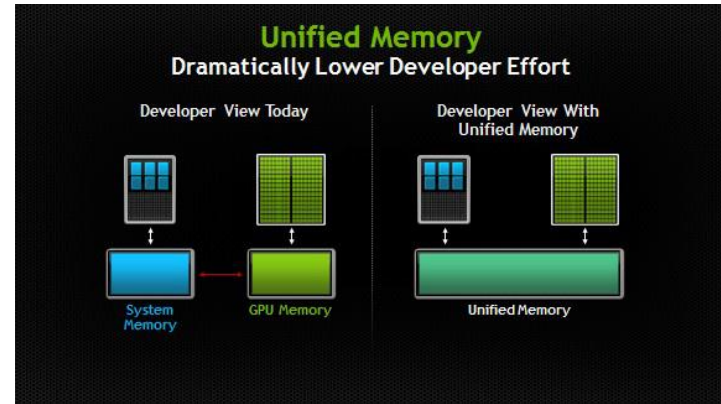
- In a typical PC or cluster node today, the memories of the CPU and GPU are physically distinct and separated by the PCI-Express bus.
- Data that is shared between the CPU and GPU must be allocated in both memories, and explicitly copied between them by the program.



Unified Memory



- Unified Memory creates a pool of managed memory that is shared between the CPU and GPU
- Managed memory is accessible to both the CPU and GPU using a single pointer
- The system **managed by runtime** automatically *migrates* data allocated in Unified Memory between host and device

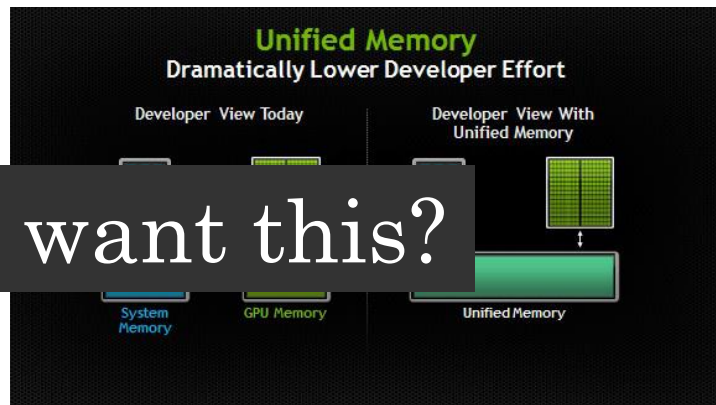


Unified Memory



- Unified Memory creates a pool of managed memory that is shared between the CPU and GPU
- Managed memory is accessible to both the CPU and GPU via a single pointer
- The system automatically *migrates* data allocated in Unified Memory between host and device

Why do we want this?



Unified Memory



- The driver handles the lookup and transfers
- Code: CPU and CUDA 6 with UM is basically the same.

CPU Code

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    data = (char *)malloc(N);  
  
    fread(data, 1, N, fp);  
  
    qsort(data, N, 1, compare);  
  
    use_data(data);  
    free(data);  
}
```

CUDA 6 Code with Unified Memory

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    cudaMallocManaged(&data, N);  
  
    fread(data, 1, N, fp);  
  
    qsort<<<...>>>(data, N, 1, compare);  
    cudaDeviceSynchronize();  
  
    use_data(data);  
    cudaFree(data);  
}
```

Unified Memory



- Notice the simplicity
 - One pointer used by the Host (CPU), and the GPU (kernels)

CPU Code

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    data = (char *)malloc(N);  
  
    fread(data, 1, N, fp);  
  
    qsort(data, N, 1, compare);  
  
    use_data(data);  
    free(data);  
}
```

CUDA 6 Code with Unified Memory

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    cudaMallocManaged(&data, N);  
  
    fread(data, 1, N, fp);  
  
    qsort<<<...>>>(data, N, 1, compare);  
    cudaDeviceSynchronize();  
  
    use_data(data);  
    cudaFree(data);  
}
```

Unified Memory



- Allows Performance through data locality
 - Data migration happens on demand between CPU and GPU.
- Despite this, tuned programs that don't use UM will perform better due to added complexity in the Driver.

CPU Code

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    data = (char *)malloc(N);  
  
    fread(data, 1, N, fp);  
  
    qsort(data, N, 1, compare);  
  
    use_data(data);  
    free(data);  
}
```

CUDA 6 Code with Unified Memory

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    cudaMallocManaged(&data, N);  
  
    fread(data, 1, N, fp);  
  
    qsort<<<...>>>(data, N, 1, compare);  
    cudaDeviceSynchronize();  
  
    use_data(data);  
    cudaFree(data);  
}
```

Streams



- Cuda stream:
 - A stream in CUDA is a sequence of operations that execute on the device in the order in which they are issued by the host code.
 - Operations from different streams can be interleaved
 - Stream IDs are used as arguments to async calls and kernel launches

```
cudaStream_t stream1, stream2, stream3, stream4 ;
cudaStreamCreate ( &stream1 ) ;
...
cudaMalloc ( &dev1, size ) ;
cudaMallocHost ( &host1, size ) ;
...
cudaMemcpyAsync ( dev1, host1, size, H2D, stream1 ) ;
kernel2 <<< grid, block, 0, stream2 >>> ( ..., dev2, ... ) ;
kernel3 <<< grid, block, 0, stream3 >>> ( ..., dev3, ... ) ;
cudaMemcpyAsync ( host4, dev4, size, D2H, stream4 ) ;
some_CPU_method () ;
...
```

// pinned memory required on host

} potentially overlapped

- Fully asynchronous / concurrent
- Data used by concurrent operations should be independent

- Cuda stream:
 - A stream in CUDA is a sequence of operations that execute on the device in the order in which they are submitted to the stream.
 - Operations from different streams can be interleaved on the device.
 - Stream IDs are used to identify streams.

`cudaMemcpyAsync`

- Asynchronous host-device memory copy returns control immediately to CPU
- Requires pinned host memory (allocated with `cudaMallocHost`)

```
kernel2 <<< grid, block, 0, stream2 >>> ( ..., dev2, ... );  
kernel3 <<< grid, block, 0, stream3 >>> ( ..., dev3, ... );  
cudaMemcpyAsync ( host4, dev4, size, D2H, stream4 );  
some_CPU_method ();  
...
```

potentially
overlapped

- Fully asynchronous / concurrent
- Data used by concurrent operations should be independent

- Synchronization
 - `cudaThreadSynchronize()`
 - Blocks until all previously issued CUDA calls from a CPU thread are complete
 - `cudaStreamSynchronize (stream):`
 - Blocks until all CUDA calls issued to given stream have completed.
 - `cudaStreamQuery (stream)`
 - Indicates whether stream is idle.
 - Doesn't block CPU thread
 - `CudaStreamWaitEvent (event)`
 - wait for event in a stream
-

- If we have multiple streams, how many different ways exist to do a complete kernel operation (H2D/Kernel/D2H)?
-

Streams

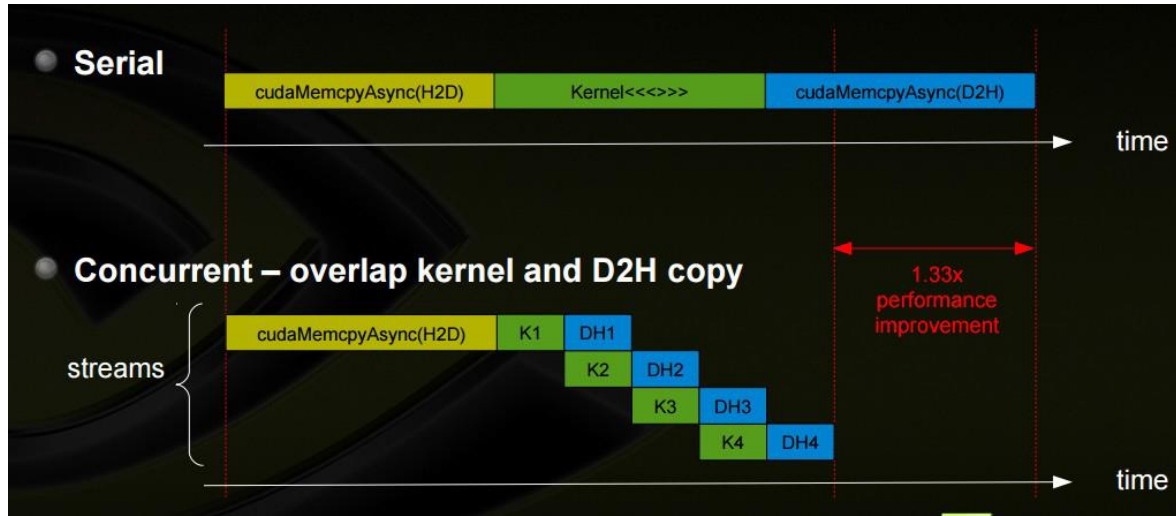


- If we have multiple streams, how many different ways exist to do a complete kernel operation (H2D/Kernel/D2H)?



Streams

- If we have multiple streams, how many different ways exist to do a complete kernel operation (H2D/Kernel/D2H)?



Streams



- If we have multiple streams, how many different ways exist to do a complete kernel operation (H2D/Kernel/D2H)?



- Example results from a Tiled DGEMM

- CPU

- 43 Gflops

- GPU

- Serial: 126 Gflops (2.9x)

- 2-way: 177 Gflops (4.1x)

- 3-way: 262 Gflops (6.1x)

- GPU + CPU

- 4-way: 282 Gflops (6.6x)



- If we have multiple streams, how many different ways exist to do a complete kernel operation (H2D/Kernel/D2H)?
 - We need to take into account how many Copy Engines are available!
 - Lets look at 2 methods on how we could do these operations.
-

- Method 1

```
for (int i = 0; i < nStreams; ++i) {  
    int offset = i * streamSize;  
    cudaMemcpyAsync(&d_a[offset], &a[offset], streamBytes, cudaMemcpyHostToDevice,  
        kernel<<<streamSize/blockSize, blockSize, 0, stream[i]>>>(d_a, offset);  
    cudaMemcpyAsync(&a[offset], &d_a[offset], streamBytes, cudaMemcpyDeviceToHost,  
        kernel<<<streamSize/blockSize, blockSize, 0, stream[i]>>>(d_a, offset);  
}
```

- Method 1
 - CHD1, K1, CDH1, CHD2, K2, CDH2, ...
- Queues:
 - Copy engine
 - CHD1, CDH1, CHD2, CDH2, ...
 - Kernel
 - K1, K2, ...

```
for (int i = 0; i < nStreams; ++i) {  
    int offset = i * streamSize;  
    cudaMemcpyAsync(&d_a[offset], &a[offset], streamBytes, cudaMemcpyHostToDevice,  
        kernel<<<streamSize/blockSize, blockSize, 0, stream[i]>>>(&d_a, offset);  
    cudaMemcpyAsync(&a[offset], &d_a[offset], streamBytes, cudaMemcpyDeviceToHost,  
        kernel<<<streamSize/blockSize, blockSize, 0, stream[i]>>>(&a, offset);  
}
```

- Method 2

```
for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    cudaMemcpyAsync(&d_a[offset], &a[offset],
                   streamBytes, cudaMemcpyHostToDevice, cudaMemcpyHostToDevice)
}

for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    kernel<<<streamSize/blockSize, blockSize, 0, stream[i]>>>(d_a, offset);
}

for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    cudaMemcpyAsync(&a[offset], &d_a[offset],
                   streamBytes, cudaMemcpyDeviceToHost, cudaMemcpyDeviceToHost)
}
```

- Method 2
 - CHD1, CHD2, ..., K1, K2, ..., CDH1, CDH2, ...

- Queues

Copy Engine

CHD1, CHD2, ... CDH1, CDH2, ...

Kernel

K1, K2, ...

```
for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    cudaMemcpyAsync(&d_a[offset], &a[offset],
                   streamBytes, cudaMemcpyHostToDevice, cudaMemcpyHostToDevice)
}

for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    kernel<<<streamSize/blockSize, blockSize, 0, stream[i]>>>(d_a, offset);
}

for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    cudaMemcpyAsync(&a[offset], &d_a[offset],
                   streamBytes, cudaMemcpyDeviceToHost, cudaMemcpyDeviceToHost)
}
```

Copy Engines




- When the device has one copy engine

Sequential Version



Asynchronous Version 1



Asynchronous Version 2



Time →

Copy Engines



- When the device has one copy engine

- METHOD 1
 - CHD1, K1,
CDH1, CHD2,
K2, CDH2, ...
- QUEUES:
 - COPY ENGINE
 - CHD1,
CDH1,
CHD2,
CDH2, ...
 - KERNEL
 - K1, K2, ...

Sequential Version



Asynchronous Version 1



Asynchronous Version 2



Time →

Copy Engines



- When the device has one copy engine

- METHOD 2
 - CHD1, CHD2, ..., K1, K2, ..., CDH1, CDH2, ...
- QUEUES
 - COPY ENGINE
 - CHD1, CHD2, ... CDH1, CDH2, ...
 - KERNEL
 - K1, K2, ...

Sequential Version



Asynchronous Version 1



Asynchronous Version 2




Copy Engines




- When the device has two copy engines (with compute capability <3.5)

Sequential Version



Asynchronous Version 1



Asynchronous Version 2




Time →

Copy Engines



- When the device has two copy engines (with compute capability <3.5)

Sequential Version




- METHOD 1
 - CHD1, K1, CDH1, CHD2, K2, CDH2, ...

QUEUES:

- COPY ENGINE
 - CHD1, CDH1, CHD2, CDH2, ...
- KERNEL
 - K1, K2, ...

Asynchronous Version 1



Asynchronous Version 2



Time →

Copy Engines

- When the device has two copy engines (with compute capability <3.5)

Sequential Version



Asynchronous Version 1



Asynchronous Version 2



Time →

- METHOD 1
 - CHD1, K1, CDH1, CHD2, K2, CDH2, ...
- QUEUES:
 - COPY ENGINE
 - CHD1, **CDH1**, CHD2, ...
 - KERNEL
 - K1, K2, ...


Added to the other Copy Engine

Copy Engines




- When the device has two copy engines (with compute capability <3.5)

Sequential Version




- METHOD 2
 - CHD1, CHD2, ..., K1, K2, ..., CDH1, CDH2,

Asynchronous Version 1



- QUEUES
 - COPY ENGINE
 - CHD1, CHD2, ... CDH1, CDH2, ...
 - KERNEL
 - K1, K2,

Asynchronous Version 2



Time →


Copy Engines

- When the device has two copy engines (with compute capability <3.5)


Sequential Version



Asynchronous Version 1



Asynchronous Version 2



Time →

- METHOD 2
 - CHD1, CHD2, ..., K1, K2, ..., CDH1, CDH2,
- QUEUES
 - COPY ENGINE
 - CHD1, CHD2, ..., CDH1, CDH2, ...
 - KERNEL
 - K1, K2,

Added to the other Copy Engine

Copy Engines




- When the device has two copy engines (with **compute capability <3.5**)

- METHOD 2
 - CHD1, CHD2, ..., K1, K2, ..., CDH1, CDH2,
- QUEUES
 - COPY ENGINE
 - CHD1, CHD2, ... CDH1, CDH2, ...
 - KERNEL
 - K1, K2,

Sequential Version



Asynchronous Version



The D2H engine won't be engaged until all kernels finish, if kernels are issued back to back.

Asynchronous Version 2



Time →

Added to the other Copy Engine

Copy Engines



- When the device has two copy engines

- METHOD 2
 - CHD1, CHD2, ..., K1, K2, ..., CDH1, CDH2,
- QUEUES
 - COPY ENGINE
 - CHD1, CHD2, ... CDH1, CDH2, ...
 - KERNEL
 - K1, K2,

Sequential Version



Asynchi



The D2H engine won't be engaged until all kernels finish, if kernels are issued back to back.

Problem if compute capability smaller <3

Time →

Added to the other Copy Engine

Copy Engines


- When the device has two copy engines (with **compute capability ≥ 3.5**)

- METHOD 1
 - CHD1, K1,
CDH1, CHD2,
K2, CDH2, ...
- QUEUES:
 - COPY ENGINE
 - CHD1, **CDH1**,
CHD2, ...
 - KERNEL
 - K1, K2, ...

Sequential Version



Asynchronous Version 1



Asynchronous Version 2



Time →

Added to the other Copy Engine

Streams again



- Now that we've copied data efficiently and overlapped with execution. Another simple example:

```
int main()
{
    const int num_streams = 8;

    cudaStream_t streams[num_streams];
    float *data[num_streams];

    for (int i = 0; i < num_streams; i++) {
        cudaStreamCreate(&streams[i]);

        cudaMalloc(&data[i], N * sizeof(float));

        // launch one worker kernel per stream
        kernel<<<1, 64, 0, streams[i]>>>(data[i], N);

        // launch a dummy kernel on the default stream
        kernel<<<1, 1>>>(0, 0);
    }

    cudaDeviceReset();

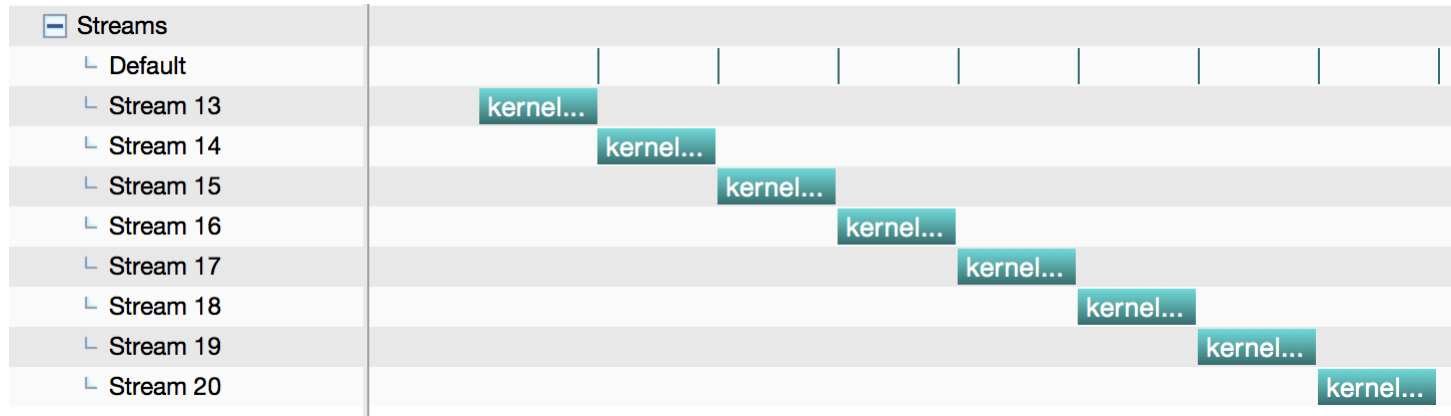
    return 0;
}
```


Streams again



- When you execute asynchronous CUDA commands without specifying a stream:
 - runtime uses the default stream.
 - Before CUDA 7, the default stream is a special stream which implicitly synchronizes with all other streams on the device.
 - CUDA 7 introduces a new option to use an independent default stream for every host thread, which avoids the serialization of the legacy default stream.
 - `--default-stream per-thread`
-

Streams again



Streams again



--default-stream per-thread

Streams		
Stream 13		kernel(float*, int)
Stream 14		
Stream 15		kernel(float*, int)
Stream 16		kernel(float*, int)
Stream 17		kernel(float*, int)
Stream 18		kernel(float*, int)
Stream 19		kernel(float*, int)
Stream 20		kernel(float*, int)
Stream 21		kernel(float*, int)

- Device-side kernel launches
 - Kepler GK110 architecture
 - Typical use cases
 - Dynamic load balancing
 - Data-dependent execution
 - Recursion
 - Library calls from kernels
 - Programmability and maintainability
-

Dynamic Parallelism

- Device-side kernel launches
 - Kepler GK110 architecture
 - Typical use cases
 - Dynamic load balancing
 - Data-dependent execution
 - Recursion
 - Library calls from kernels
 - Programmability and maintainability



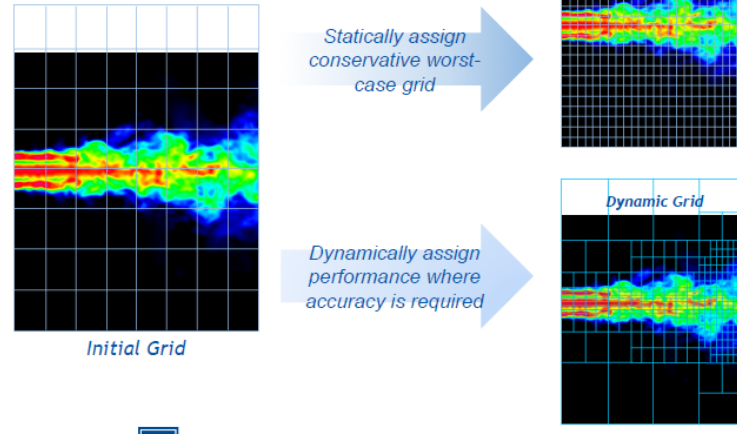
Fermi: Only CPU can generate GPU work



Kepler: GPU can generate work for itself

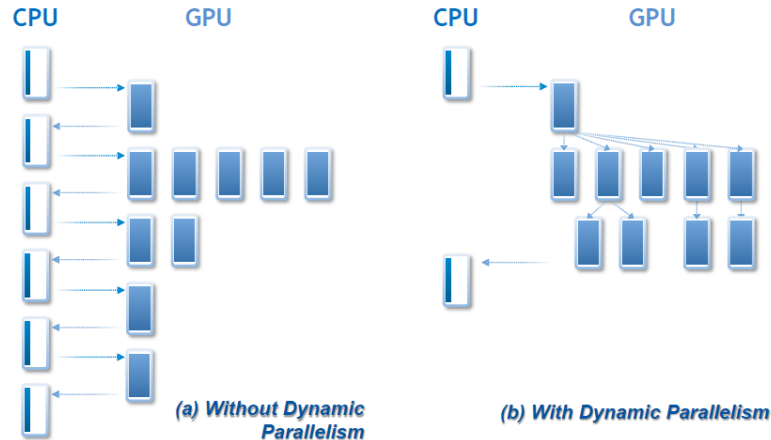
Dynamic Parallelism

- Fixed grid vs dynamic grid for a turbulence simulation mode



Dynamic Parallelism

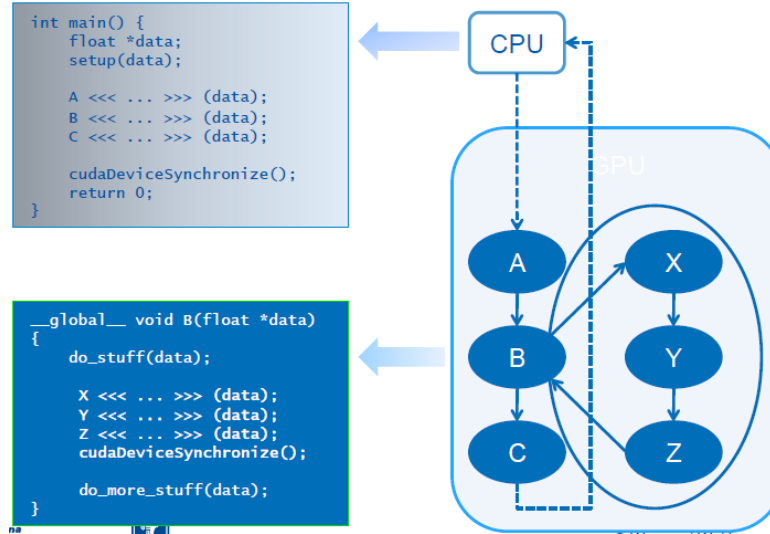
- CPU-GPU without and with dynamic parallelism



Dynamic Parallelism



- Nested dependencies



- Syntax

```
Kernel_name <<<Dg, Db, Ns, S>>> ([kernel arguments]);
```

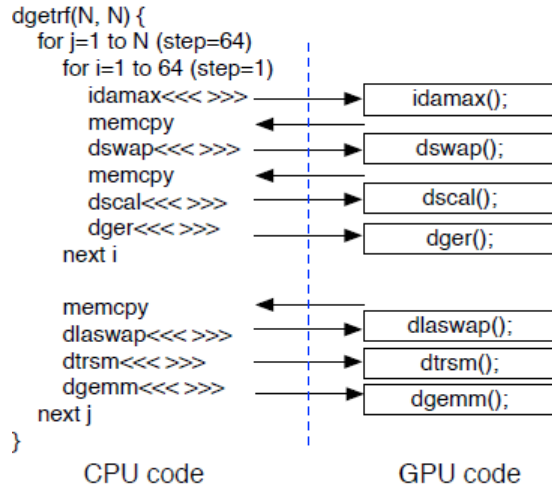
- Dg is of type dim3 and specifies the dimensions and size of the grid
 - Db is of type dim3 and specifies the dimensions and size of each thread block
 - Ns is of type size_t and specifies the number of bytes of shared memory that is dynamically allocated per thread block for this call
 - S is of type cudaStream_t and specifies the stream associated with this call
-

Dynamic Parallelism

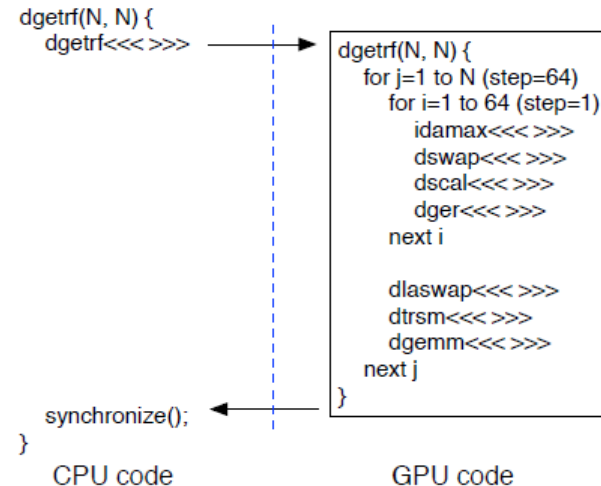


- Example

LU decomposition (Fermi)



LU decomposition (Kepler)



Dynamic Parallelism




- Synchronization

- Parent to child: memory consistency

- Child to parent: after

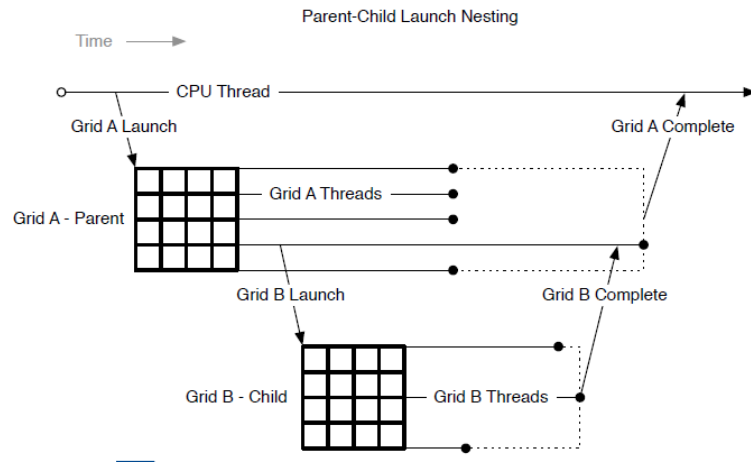
- `cudaDeviceSynchronize()`



Dynamic Parallelism



- Memory Model
 - Child sees parent state at time of launch
 - Parent sees child writes after sync
 - Constants are immutable
 - Local and shared memory are private



Dynamic Parallelism




- Launch pool size
 - Fixed-size pool: default 2048
 - Variable-size pool

Before CUDA 6.0



Since CUDA 6.0

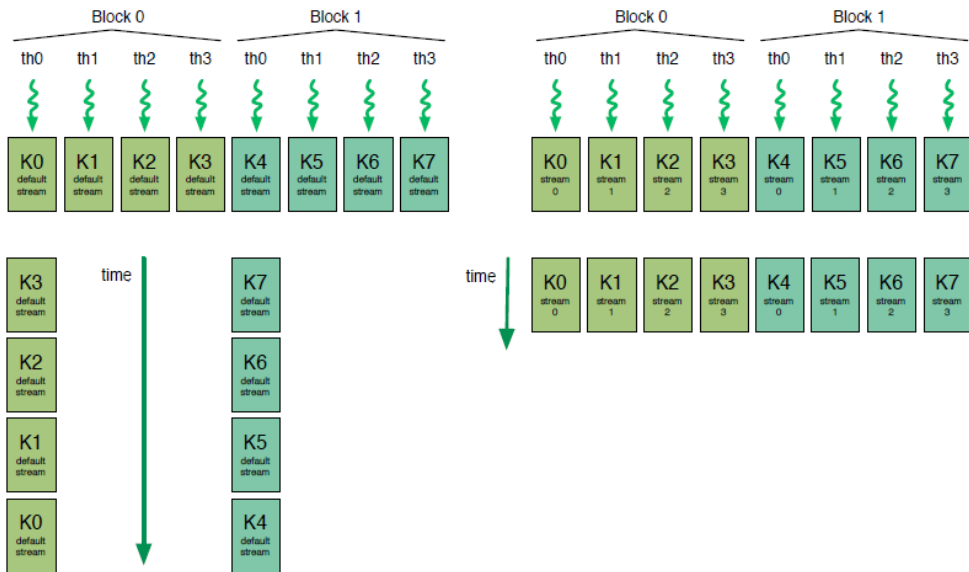


Dynamic Parallelism





- Streams

To guarantee concurrency



- Level Set Algorithm




- Level Set Methods

- Used to detect the border of an object on an image
- They use partial differential equations to evolve the curve in the image
- Positive = object
- Negative = not object

Dynamic Parallelism



- Level Set Algorithm Implementation



```
checkCuda(cudaMemcpy(gpu.intensity,
intensity,
gpu.size*sizeof(int),
cudaMemcpyHostToDevice));

checkCuda(cudaMemcpy(gpu.labels,
labels,
gpu.size*sizeof(int),
cudaMemcpyHostToDevice));

checkCuda(cudaDeviceSynchronize());

#ifdef(CUDA_TIMING)
float Ktime;
TIMER_CREATE(Ktime);
TIMER_START(Ktime);
#endif

#ifdef(VERBOSE)
printf("Running algorithm on GPU.\n");
#endif

// Launch kernel to begin image segmentation
evolveContour<<<1, numLabels>>>(gpu.intensity,
gpu.labels,
gpu.phi,
gpu.phiOut,
gridXSize,
gridYSize,
gpu.targetLabels,
gpu.lowerIntensityBounds,
gpu.upperIntensityBounds,
max_iterations,
gpu.globalBlockIndicator,
gpu.globalFinishedVariable,
gpu.totalIterations);


checkCuda(cudaDeviceSynchronize());

#ifdef(CUDA_TIMING)
TIMER_END(Ktime);
printf("Kernel Execution Time: %f ms\n", Ktime);
#endif
```


Dynamic Parallelism



- Level Set Algorithm Implementation



```
checkCuda(cudaMemcpy(gpu.intensity,
intensity,
gpu.size*sizeof(int),
cudaMemcpyHostToDevice));
```

```
checkCuda(cudaMemcpy(gpu.labels,
labels,
gpu.size*sizeof(int),
cudaMemcpyHostToDevice));
```

```
checkCuda(cudaDeviceSynchronize());
```

```
#if defined(CUDA_TIMING)
float Ktime;
TIMER_CREATE(Ktime);
TIMER_START(Ktime);
#endif
```

```
#if defined(VERBOSE)
printf("Running algorithm on GPU.\n");
#endif
```

```
// Launch kernel to begin image segmentation
evolveContour<<<1, numLabels>>>>(gpu.intensity,
gpu.labels,
gpu.phi,
gpu.phiOut,
gridXSize,
gridYSize,
gpu.targetLabels,
gpu.lowerIntensityBounds,
gpu.upperIntensityBounds,
max_iterations,
gpu.globalBlockIndicator,
gpu.globalFinishedVariable,
gpu.totalIterations);
```



```
evolveContour(gpu.intensity, gpu.labels, gpu.phi, gpu.phiOut, gridXSize, gridYSize, gpu.targetLabels, gpu.lowerIntensityBounds, gpu.upperIntensityBounds, max_iterations, gpu.globalBlockIndicator, gpu.globalFinishedVariable, gpu.totalIterations);
```

```
#if defined(CUDA_TIMING)
TIMER_END(Ktime);
printf("Kernel Execution Time: %f ms\n", Ktime);
#endif
```

Dynamic Parallelism



- Level Set Algorithm Implementation



```

_global_ void evolveContour(unsigned int* intensity,
                           unsigned int* labels,
                           signed int* phi,
                           signed int* phiOut,
                           int gridSize,
                           int gridSize,
                           int* targetLabels,
                           int* lowerIntensityBounds,
                           int* upperIntensityBounds,
                           int max_iterations,
                           int* globalBlockIndicator,
                           int* globalFinishedVariable,
                           int* totalIterations ) {

    int tid = threadIdx.x;

    // Setting up streams for
    cudaStream_t stream;
    cudaStreamCreateWithFlags (&stream, cudaStreamNonBlocking);

    // Total iterations
    totalIterations = &totalIterations[tid];

    // Size in ints
    int size = (gridSize*BLOCK_TILE_SIZE)*(gridSize*BLOCK_TILE_SIZE);

    // New phi pointer for each label.
    phi = &phi[tid*size];
    phiOut = &phiOut[tid*size];

    globalBlockIndicator = &globalBlockIndicator[tid*gridSize*gridSize];

    // Global synchronization variable
    globalFinishedVariable = &globalFinishedVariable[tid];

    dim3 dimGrid(gridSize, gridSize);
    dim3 dimBlock(BLOCK_TILE_SIZE, BLOCK_TILE_SIZE);

    // Initialize phi array
    lssStep1<<<dimGrid, dimBlock, 0, stream>>>(intensity,
                                                labels,
                                                phi,
                                                targetLabels[tid],
                                                lowerIntensityBounds[tid],
                                                upperIntensityBounds[tid],
                                                globalBlockIndicator);

    int iterations = 0;
    do {
        iterations++;
        lssStep2<<<dimGrid, dimBlock, 0, stream>>>(phi,
                                                    globalBlockIndicator,
                                                    globalFinishedVariable );
        cudaDeviceSynchronize();
    } while (atomicExch(globalFinishedVariable,0) && (iterations < max_iterations));




    lssStep3<<<dimGrid, dimBlock, 0, stream>>>(phi,
                                                phiOut);

    *totalIterations = iterations;
}
    
```

Dynamic Parallelism



- Level Set Algorithm Implementation



```
_global_ void evolveContour(unsigned int* intensity,
                           unsigned int* labels,
                           signed int* phi,
                           signed int* phiOut,
                           int gridSize,
                           int gridSize,
                           int* targetLabels,
                           int* lowerIntensityBounds,
                           int* upperIntensityBounds,
                           int max_iterations,
                           int* globalBlockIndicator,
                           int* globalFinishedVariable,
                           int* totalIterations ) {

    int tid = threadIdx.x;

    // Setting up streams for
    cudaStream_t stream;
    cudaStreamCreateWithFlags (&stream, cudaStreamNonBlocking);

    // Total iterations
    totalIterations = &totalIterations[tid];

    // Size in ints
    int size = (gridSize*BLOCK_TILE_SIZE)*(gridSize*BLOCK_TILE_SIZE);

    // New phi pointer for each label.
    phi = &phi[tid*size];
    phiOut = &phiOut[tid*size];

    globalBlockIndicator = &globalBlockIndicator[tid*gridSize*gridSize];

    // Global synchronization variable
    globalFinishedVariable = &globalFinishedVariable[tid];

    dim3 dimGrid(gridSize, gridSize);
    dim3 dimBlock(BLOCK_TILE_SIZE, BLOCK_TILE_SIZE);

    // Initialize phi array
    lssStep1<<<dimGrid, dimBlock, 0, stream>>>(intensity,
                                                labels,
                                                phi,
                                                targetLabels[tid],
                                                lowerIntensityBounds[tid],
                                                upperIntensityBounds[tid],
                                                globalBlockIndicator);

    int iterations = 0;
    do {
        iterations++;
        lssStep2<<<dimGrid, dimBlock, 0, stream>>>(phi,
                                                    globalBlockIndicator,
                                                    globalFinishedVariable );
        cudaDeviceSynchronize();
    } while (atomicExch(globalFinishedVariable,0) && (iterations < max_iterations));
    lssStep3<<<dimGrid, dimBlock, 0, stream>>>(phi,
                                                phiOut);

    *totalIterations = iterations;
}
```

- To understand recursive usage
 - Classic Fibonacci series: 0, 1, 1, 2, 3, 5...
 - CPU Recursive
 - GPU non recursive nonDP
 - GPU recursive DP

Fibonacci – CPU Recursion



```
int fib(int n){  
    if (n == 0 || n == 1){  
        return n;  
    }  
    else{  
        return fib(n-1) + fib(n-2);  
    }  
}
```

Fibonacci – GPU Basic Kernel



```
__global__ void fib_kernel_plain(int n, long int* vFib){  
    int tid = threadIdx.x + blockDim.x * blockIdx.x;  
    if (tid > n/32)  
        return;  
  
    if (n == 0 || n == 1){  
        return;  
    }  
    for(int i=tid*32 + 2; i <= n && i < tid*32 + 32; i++){  
        vFib[i] = vFib[i-1] + vFib[i-2];  
    }  
}
```


Fibonacci – GPU Recursion Kernel

```
__global__ void fib_kernel_par_rec(int n, unsigned long int* vFib){  
    if (n == 0 || n == 1)  
        return;  
  
    fib_kernel_par_rec<<<1, 1>>>(n-2, vFib);  
    fib_kernel_par_rec<<<1, 1>>>(n-1, vFib);  
    cudaDeviceSynchronize();  
    vFib[n] = vFib[n-1] + vFib[n-2];  
  
}
```

Results




- Performance CPU vs GPU non DP vs GPU DP




Results

- Performance CPU vs GPU non DP vs GPU DP



Results

- Performance CPU vs GPU non DP vs GPU DP



What is and what is not dynamic parallelism



- CDP ensures better work balance, and offers advantages in terms of programmability
 - However, launching grids with a very small number of threads could lead to severe underutilization of the GPU resources
 - A general recommendation
 - Child grids with a large number of thread blocks,
 - Or at least thread blocks with hundreds of threads, if the number of blocks is small
 - Nested parallelism for tree processing
 - Thick tree nodes (each node deploys many threads) work well
 - And/or when branch degree is large (each parent node has many children)
 - As the nesting depth is limited in hardware, only relatively shallow trees can be implemented efficiently.
-