



GPU Programming

(in CUDA)

Summer 2020

Designed by Julian Gutierrez, Presented by Nicolas Agostini

Session 11





- Improving Host to Device Memory Transfers
 - Pinned Memory
 - Unified Memory
- Concurrency
- Dynamic Parallelism

An abstract graphic on a dark blue background with a light blue grid. It features a glowing blue node at the center, with a network of white lines and points extending from it. A series of white lines, resembling a circuit board, extends from the node towards the bottom right corner.

Improving Host to Device Memory Transfers

Memory Transfers



- We want to minimize the amount of time it takes to transfer data between host and device.
- Device to host memory bandwidth is much lower than device to device bandwidth

Memory Transfers



- We want to minimize the amount of time it takes to transfer data between host and device.
 - How can we do this?
 - Reduce the amount of data we have to transfer! (if possible)
 - Achieve highest memory bandwidth between host and device. (Possible when using page-locked (or “pinned”) memory, we are going to talk about it in this session)
 - Batching many small transfers into one larger transfer performs much better because it eliminates most of the per-transfer overhead
 - Data transfers between the host and device can sometimes be overlapped with kernel execution and other data transfers (we will talk about this in this session as well)
-


- Remember
 - When deciding whether to run on the GPU or on the CPU, we should consider the kernel execution time and the time it takes to copy data back and forth.
 - We need to consider the cost of moving data across the PCI-e bus.
-

Pinned Memory (Page-Locked Data Transfers)




- Host (CPU) data allocations are pageable by default.
 - The GPU cannot access data directly from pageable host memory.
 - Due to this, the Cuda driver must:
 - Allocate a temporary page-locked block, or “pinned”
 - Copy host data to the pinned block
 - Transfer from pinned to device
 - Delete pinned block

Pageable Data Transfer



Pinned Data Transfer




Pinned Memory (Page-Locked Data Transfers)




- `cudaMallocHost()`
 - Prevents OS from paging host memory
 - Allows PCI-e DMA to run at full speed
- `cudaFreeHost()`
- Allocations can fail so you need to check errors.

Pageable Data Transfer



`malloc`

Pinned Data Transfer



`cudaMallocHost`

```
// allocate and initialize
h_aPageable = (float*)malloc(bytes); // host pageable
h_bPageable = (float*)malloc(bytes); // host pageable
checkCuda( cudaMallocHost((void**)&h_aPinned, bytes) ); //
host pinned
checkCuda( cudaMallocHost((void**)&h_bPinned, bytes) ); //
host pinned
checkCuda( cudaMalloc((void**)&d_a, bytes) ); // device
```

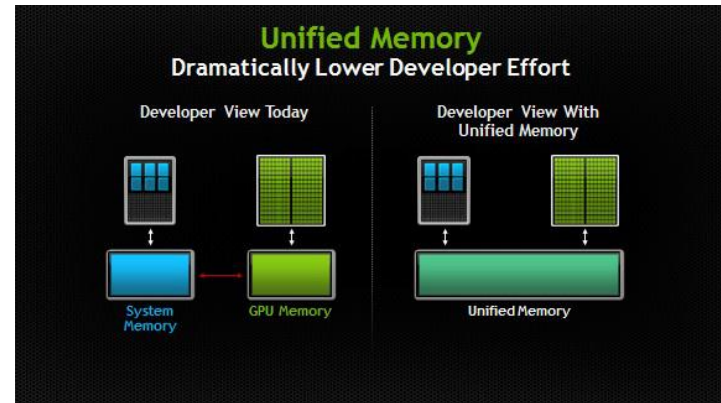

Pinned Memory (Page-Locked Data Transfers)

- You should not over-allocate pinned memory. Doing so can reduce overall system performance because it reduces the amount of physical memory available to the operating system and other programs.
 - Transfers between the host and device are the slowest link of data movement involved in GPU computing, so you should take care to minimize transfers.
-

Unified Memory



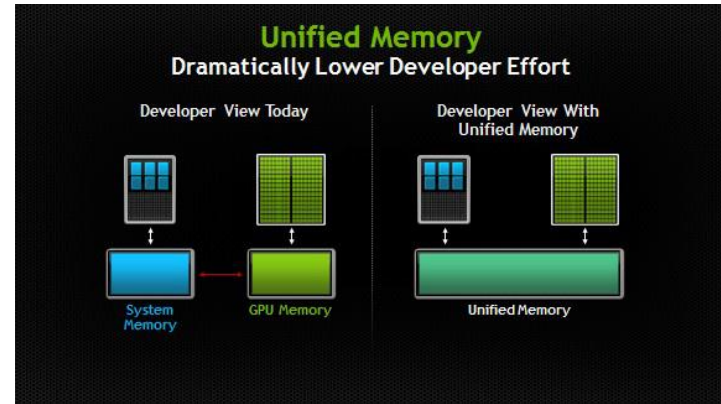
- In a typical PC or cluster node today, the memories of the CPU and GPU are physically distinct and separated by the PCI-Express bus.
- Data that is shared between the CPU and GPU must be allocated in both memories, and explicitly copied between them by the program.



Unified Memory



- Unified Memory creates a pool of managed memory that is shared between the CPU and GPU
- Managed memory is accessible to both the CPU and GPU using a single pointer
- The system automatically *migrates* data allocated in Unified Memory between host and device

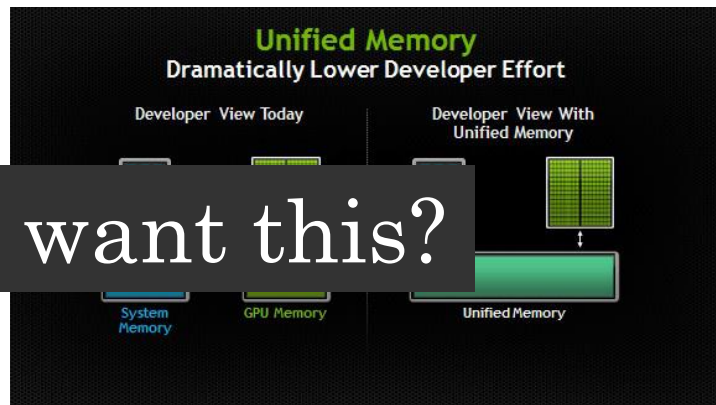


Unified Memory



- Unified Memory creates a pool of managed memory that is shared between the CPU and GPU
- Managed memory is accessible to both the CPU and GPU via a single pointer
- The system automatically *migrates* data allocated in Unified Memory between host and device

Why do we want this?



Unified Memory



- The driver handles the lookup and transfers
- Code: CPU and CUDA 6 with UM is basically the same.

CPU Code

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    data = (char *)malloc(N);  
  
    fread(data, 1, N, fp);  
  
    qsort(data, N, 1, compare);  
  
    use_data(data);  
    free(data);  
}
```

CUDA 6 Code with Unified Memory

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    cudaMallocManaged(&data, N);  
  
    fread(data, 1, N, fp);  
  
    qsort<<<...>>>(data, N, 1, compare);  
    cudaDeviceSynchronize();  
  
    use_data(data);  
    cudaFree(data);  
}
```

Unified Memory



- Notice the simplicity
 - One pointer used by the Host (CPU), and the GPU (kernels)

CPU Code

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    data = (char *)malloc(N);  
  
    fread(data, 1, N, fp);  
  
    qsort(data, N, 1, compare);  
  
    use_data(data);  
    free(data);  
}
```

CUDA 6 Code with Unified Memory

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    cudaMallocManaged(&data, N);  
  
    fread(data, 1, N, fp);  
  
    qsort<<<...>>>(data, N, 1, compare);  
    cudaDeviceSynchronize();  
  
    use_data(data);  
    cudaFree(data);  
}
```

Unified Memory



- Allows Performance through data locality
 - Data migration happens on demand between CPU and GPU.
- Despite this, tuned programs that don't use UM will perform better due to added complexity in the Driver.

CPU Code

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    data = (char *)malloc(N);  
  
    fread(data, 1, N, fp);  
  
    qsort(data, N, 1, compare);  
  
    use_data(data);  
    free(data);  
}
```

CUDA 6 Code with Unified Memory

```
void sortfile(FILE *fp, int N) {  
    char *data;  
    cudaMallocManaged(&data, N);  
  
    fread(data, 1, N, fp);  
  
    qsort<<<...>>>(data, N, 1, compare);  
    cudaDeviceSynchronize();  
  
    use_data(data);  
    cudaFree(data);  
}
```

Streams



- **Cuda stream:**
 - A stream in CUDA is a sequence of operations that execute on the device in the order in which they are issued by the host code.
 - Operations from different streams can be interleaved
 - Stream IDs are used as arguments to async calls and kernel launches

```
cudaStream_t stream1, stream2, stream3, stream4 ;
cudaStreamCreate ( &stream1 ) ;
...
cudaMalloc ( &dev1, size ) ;
cudaMallocHost ( &host1, size ) ;
...
cudaMemcpyAsync ( dev1, host1, size, H2D, stream1 ) ;
kernel2 <<< grid, block, 0, stream2 >>> ( ..., dev2, ... ) ;
kernel3 <<< grid, block, 0, stream3 >>> ( ..., dev3, ... ) ;
cudaMemcpyAsync ( host4, dev4, size, D2H, stream4 ) ;
some_CPU_method () ;
...
```

// pinned memory required on host

} potentially overlapped

- **Fully asynchronous / concurrent**
- **Data used by concurrent operations should be independent**

Streams



- Cuda stream:
 - A stream in CUDA is a sequence of operations that execute on the device in the order in which they are submitted.
 - Operations from different streams can be interleaved on the device.
 - Stream IDs are used to identify the stream of an operation.

cudaMemcpyAsync

- Asynchronous host-device memory copy returns control immediately to CPU
- Requires pinned host memory (allocated with cudaMallocHost)

```
kernel2 <<< grid, block, 0, stream2 >>> ( ..., dev2, ... );  
kernel3 <<< grid, block, 0, stream3 >>> ( ..., dev3, ... );  
cudaMemcpyAsync ( host4, dev4, size, D2H, stream4 );  
some_CPU_method ();  
...
```

potentially
overlapped

- Fully asynchronous / concurrent
- Data used by concurrent operations should be independent

- Synchronization
 - `cudaThreadSynchronize()`
 - Blocks until all previously issued CUDA calls from a CPU thread are complete
 - `cudaStreamSynchronize (stream):`
 - Blocks until all CUDA calls issued to given stream have completed.
 - `cudaStreamQuery (stream)`
 - Indicates whether stream is idle.
 - Doesn't block CPU thread
 - `CudaStreamWaitEvent (event)`
 - wait for event in a stream
-

- If we have multiple streams, how many different ways exist to do a complete kernel operation (H2D/Kernel/D2H)?
-

Streams

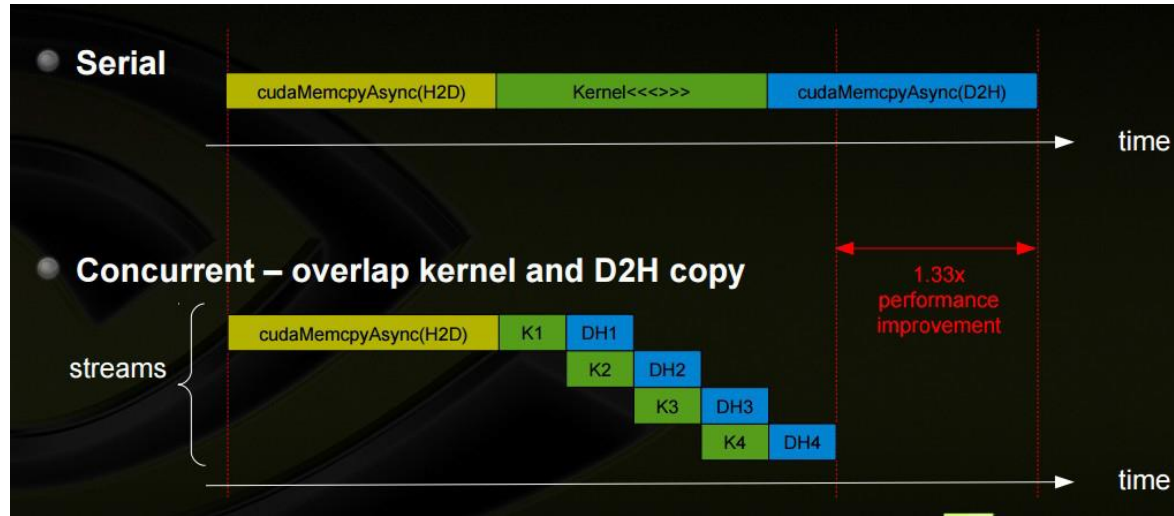


- If we have multiple streams, how many different ways exist to do a complete kernel operation (H2D/Kernel/D2H)?



Streams

- If we have multiple streams, how many different ways exist to do a complete kernel operation (H2D/Kernel/D2H)?



Streams

- If we have multiple streams, how many different ways exist to do a complete kernel operation (H2D/Kernel/D2H)?



- Example results from a Tiled DGEMM

- CPU

- 43 Gflops

- GPU


- Serial: 126 Gflops (2.9x)

- 2-way: 177 Gflops (4.1x)

- 3-way: 262 Gflops (6.1x)

- GPU + CPU

- 4-way: 282 Gflops (6.6x)



- If we have multiple streams, how many different ways exist to do a complete kernel operation (H2D/Kernel/D2H)?
 - We need to take into account how many Copy Engines are available!
 - Lets look at 2 methods on how we could do these operations.
-

- Method 1

```
for (int i = 0; i < nStreams; ++i) {  
    int offset = i * streamSize;  
    cudaMemcpyAsync(&d_a[offset], &a[offset], streamBytes, cudaMemcpyHostToDevice,  
        kernel<<<streamSize/blockSize, blockSize, 0, stream[i]>>>(d_a, offset);  
    cudaMemcpyAsync(&a[offset], &d_a[offset], streamBytes, cudaMemcpyDeviceToHost,  
        kernel<<<streamSize/blockSize, blockSize, 0, stream[i]>>>(d_a, offset);  
}
```

- Method 1
 - CHD1, K1, CDH1, CHD2, K2, CDH2, ...
- Queues:
 - Copy engine
 - CHD1, CDH1, CHD2, CDH2, ...
 - Kernel
 - K1, K2, ...

```
for (int i = 0; i < nStreams; ++i) {  
    int offset = i * streamSize;  
    cudaMemcpyAsync(&d_a[offset], &a[offset], streamBytes, cudaMemcpyHostToDevice,  
        kernel<<<streamSize/blockSize, blockSize, 0, stream[i]>>>(&d_a, offset);  
    cudaMemcpyAsync(&a[offset], &d_a[offset], streamBytes, cudaMemcpyDeviceToHost,  
        kernel<<<streamSize/blockSize, blockSize, 0, stream[i]>>>(&a, offset);  
}
```

- Method 2

```
for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    cudaMemcpyAsync(&d_a[offset], &a[offset],
                   streamBytes, cudaMemcpyHostToDevice, cudaMemcpyHostToDevice)
}

for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    kernel<<<streamSize/blockSize, blockSize, 0, stream[i]>>>(d_a, offset);
}

for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    cudaMemcpyAsync(&a[offset], &d_a[offset],
                   streamBytes, cudaMemcpyDeviceToHost, cudaMemcpyDeviceToHost)
}
```

Copy Engines



- Method 2
 - CHD1, CHD2, ..., K1, K2, ..., CDH1, CDH2, ...

- Queues

Copy Engine

CHD1, CHD2, ... CDH1, CDH2, ...

Kernel

K1, K2, ...

```
for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    cudaMemcpyAsync(&d_a[offset], &a[offset],
                   streamBytes, cudaMemcpyHostToDevice, cudaMemcopyHostToDevice)
}

for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    kernel<<<streamSize/blockSize, blockSize, 0, stream[i]>>>(d_a, offset);
}

for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    cudaMemcpyAsync(&a[offset], &d_a[offset],
                   streamBytes, cudaMemcpyDeviceToHost, cudaMemcopyDeviceToHost)
}
```

Copy Engines




- When the device has one copy engine

Sequential Version



Asynchronous Version 1



Asynchronous Version 2



Time →

Copy Engines



- When the device has one copy engine

- METHOD 1
 - CHD1, K1,
CDH1, CHD2,
K2, CDH2, ...
- QUEUES:
 - COPY ENGINE
 - CHD1,
CDH1,
CHD2,
CDH2, ...
 - KERNEL
 - K1, K2, ...


Sequential Version



Asynchronous Version 1



Asynchronous Version 2



Time →

Copy Engines



- When the device has one copy engine

- METHOD 2
 - CHD1, CHD2, ..., K1, K2, ..., CDH1, CDH2, ...
- QUEUES
 - COPY ENGINE
 - CHD1, CHD2, ... CDH1, CDH2, ...
 - KERNEL
 - K1, K2, ...

Sequential Version



Asynchronous Version 1



Asynchronous Version 2




Copy Engines




- When the device has two copy engines (with compute capability < 3.5)

Sequential Version



Asynchronous Version 1



Asynchronous Version 2




Time →

Copy Engines



- When the device has two copy engines (with compute capability <3.5)

Sequential Version




- METHOD 1
 - CHD1, K1, CDH1, CHD2, K2, CDH2, ...

QUEUES:

- COPY ENGINE
 - CHD1, CDH1, CHD2, CDH2, ...
- KERNEL
 - K1, K2, ...

Asynchronous Version 1



Asynchronous Version 2




Time →


Copy Engines

- When the device has two copy engines (with compute capability <3.5)

Sequential Version



Asynchronous Version 1



Asynchronous Version 2



Time →

- METHOD 1
 - CHD1, K1, CDH1, CHD2, K2, CDH2, ...
- QUEUES:
 - COPY ENGINE
 - CHD1, **CDH1**, CHD2, ...
 - KERNEL
 - K1, K2, ...

Added to the other Copy Engine

Copy Engines




- When the device has two copy engines (with compute capability <3.5)


Sequential Version



Asynchronous Version 1



Asynchronous Version 2



Time →

- METHOD 2
 - CHD1, CHD2, ..., K1, K2, ..., CDH1, CDH2,
- QUEUES
 - COPY ENGINE
 - CHD1, CHD2, ... CDH1, CDH2, ...
 - KERNEL
 - K1, K2,


Copy Engines

- When the device has two copy engines (with compute capability <3.5)


Sequential Version



Asynchronous Version 1



Asynchronous Version 2



Time →

- METHOD 2
 - CHD1, CHD2, ..., K1, K2, ..., CDH1, CDH2,
- QUEUES
 - COPY ENGINE
 - CHD1, CHD2, ..., CDH1, CDH2, ...
 - KERNEL
 - K1, K2,

Added to the other Copy Engine

Copy Engines


- When the device has two copy engines (with **compute capability <3.5**)

- METHOD 2
 - CHD1, CHD2, ..., K1, K2, ..., CDH1, CDH2,
- QUEUES
 - COPY ENGINE
 - CHD1, CHD2, ... CDH1, CDH2, ...
 - KERNEL
 - K1, K2,

Sequential Version



Asynchroni



The D2H engine won't be engaged until all kernels finish, if kernels are issued back to back.

Asynchronous Version 2



Time →


Added to the other Copy Engine

Copy Engines



- When the device has two copy engines

- METHOD 2
 - CHD1, CHD2, ..., K1, K2, ..., CDH1, CDH2,
- QUEUES
 - COPY ENGINE
 - CHD1, CHD2, ... CDH1, CDH2, ...
 - KERNEL
 - K1, K2,




Added to the other Copy Engine

Sequential Version



Asynchi



The D2H engine won't be engaged until all kernels finish, if kernels are issued back to back.

Problem if compute capability smaller < 3



Copy Engines

- When the device has two copy engines (with **compute capability ≥ 3.5**)

- METHOD 1

- CHD1, K1,
CDH1, CHD2,
K2, CDH2, ...

- QUEUES:

- COPY ENGINE

- CHD1, **CDH1**,
CHD2, ...


- KERNEL

- K1, K2, ...


Sequential Version



Asynchronous Version 1



Asynchronous Version 2



Time →

Added to the other Copy Engine

Streams again



- Now that we've copied data efficiently and overlapped with execution. Another simple example:

```
int main()
{
    const int num_streams = 8;

    cudaStream_t streams[num_streams];
    float *data[num_streams];

    for (int i = 0; i < num_streams; i++) {
        cudaStreamCreate(&streams[i]);

        cudaMalloc(&data[i], N * sizeof(float));

        // launch one worker kernel per stream
        kernel<<<1, 64, 0, streams[i]>>>(data[i], N);

        // launch a dummy kernel on the default stream
        kernel<<<1, 1>>>(0, 0);
    }

    cudaDeviceReset();

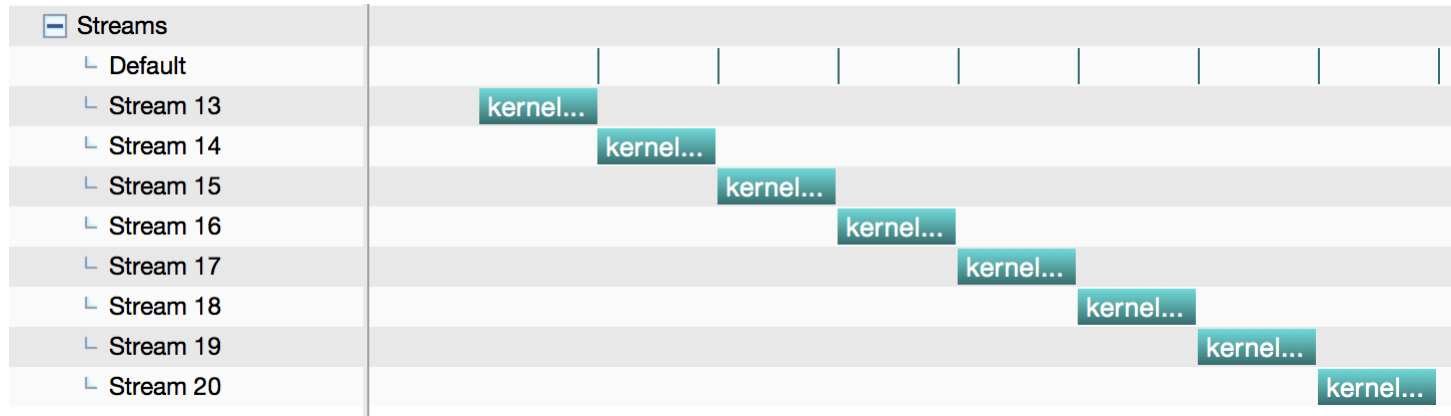
    return 0;
}
```


Streams again



- When you execute asynchronous CUDA commands without specifying a stream:
 - runtime uses the default stream.
 - Before CUDA 7, the default stream is a special stream which implicitly synchronizes with all other streams on the device.
 - CUDA 7 introduces a new option to use an independent default stream for every host thread, which avoids the serialization of the legacy default stream.
 - `--default-stream per-thread`
-

Streams again



Streams again



--default-stream per-thread

Streams		
Stream 13		kernel(float*, int)
Stream 14		
Stream 15		kernel(float*, int)
Stream 16		kernel(float*, int)
Stream 17		kernel(float*, int)
Stream 18		kernel(float*, int)
Stream 19		kernel(float*, int)
Stream 20		kernel(float*, int)
Stream 21		kernel(float*, int)

- Device-side kernel launches
 - Kepler GK110 architecture
 - Typical use cases
 - Dynamic load balancing
 - Data-dependent execution
 - Recursion
 - Library calls from kernels
 - Programmability and maintainability
-

Dynamic Parallelism

- Device-side kernel launches
 - Kepler GK110 architecture
 - Typical use cases
 - Dynamic load balancing
 - Data-dependent execution
 - Recursion
 - Library calls from kernels
 - Programmability and maintainability

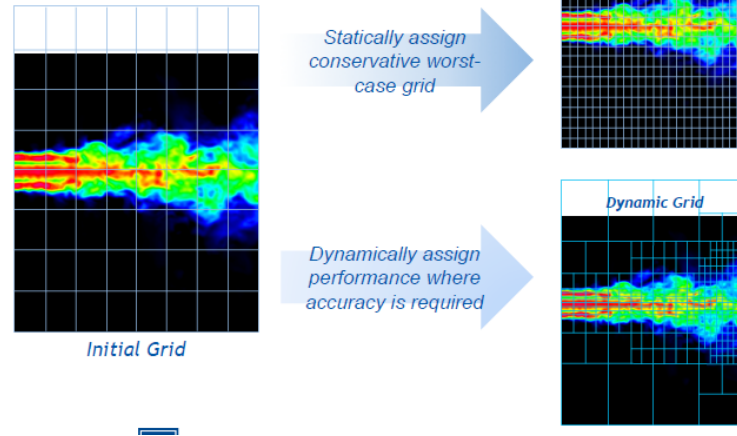


Fermi: Only CPU can generate GPU work



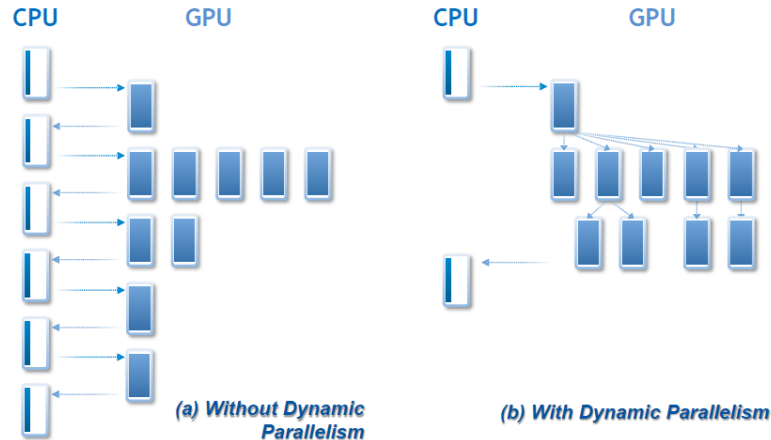
Kepler: GPU can generate work for itself

- Fixed grid vs dynamic grid for a turbulence simulation mode



Dynamic Parallelism

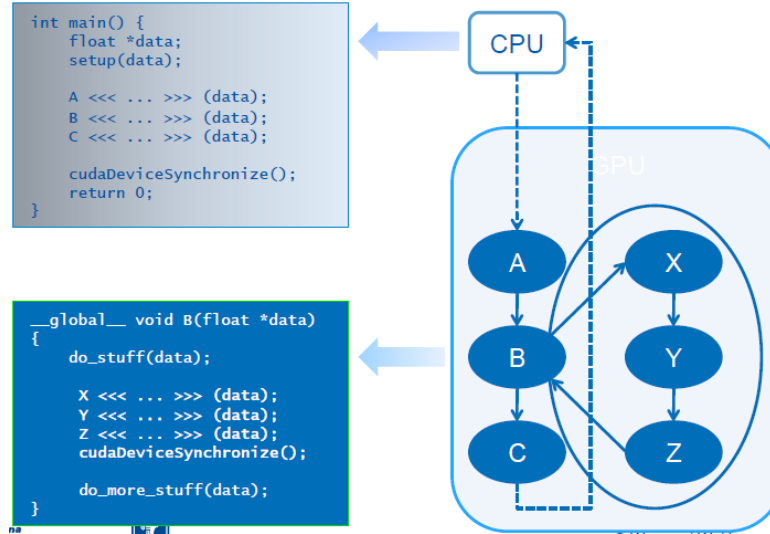
- CPU-GPU without and with dynamic parallelism



Dynamic Parallelism



- Nested dependencies



- Syntax

```
Kernel_name <<<Dg, Db, Ns, S>>> ([kernel arguments]);
```

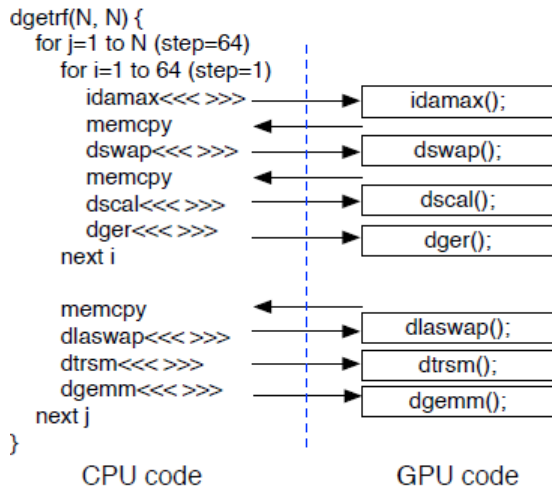
- Dg is of type dim3 and specifies the dimensions and size of the grid
 - Db is of type dim3 and specifies the dimensions and size of each thread block
 - Ns is of type size_t and specifies the number of bytes of shared memory that is dynamically allocated per thread block for this call
 - S is of type cudaStream_t and specifies the stream associated with this call
-

Dynamic Parallelism

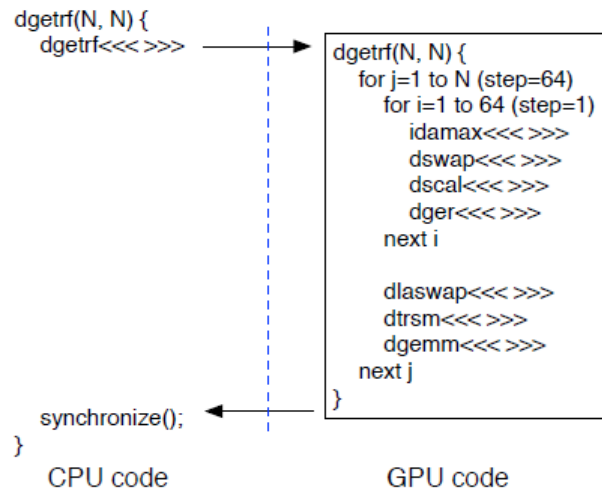


- Example

LU decomposition (Fermi)



LU decomposition (Kepler)



Dynamic Parallelism




- Synchronization

- Parent to child: memory consistency

- Child to parent: after

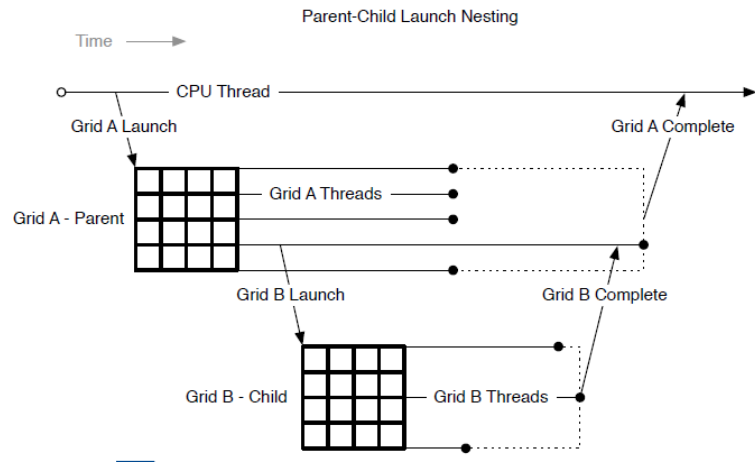
- `cudaDeviceSynchronize()`



Dynamic Parallelism



- Memory Model
 - Child sees parent state at time of launch
 - Parent sees child writes after sync
 - Constants are immutable
 - Local and shared memory are private



Dynamic Parallelism




- Launch pool size
 - Fixed-size pool: default 2048
 - Variable-size pool

Before CUDA 6.0



Since CUDA 6.0

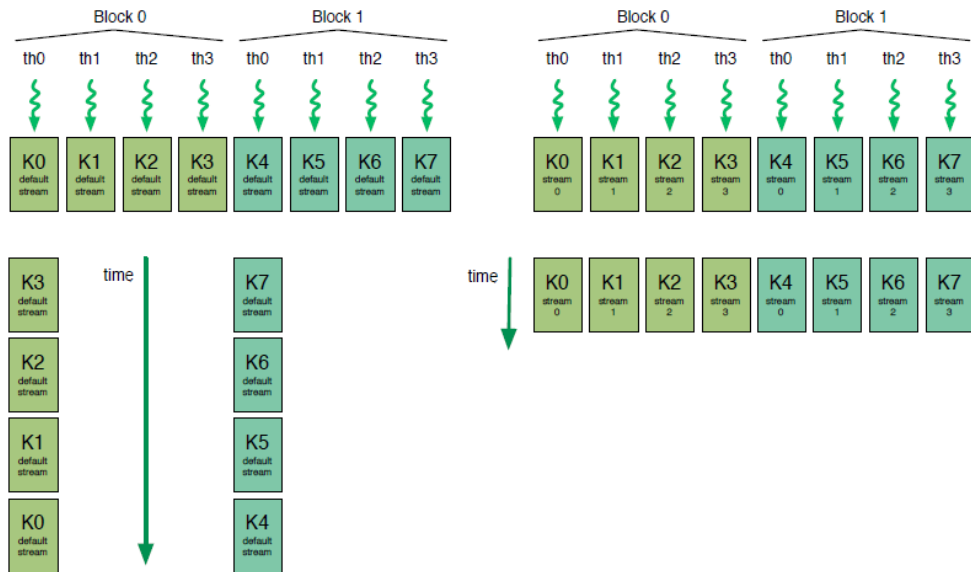


Dynamic Parallelism





- Streams

To guarantee concurrency



- Level Set Algorithm




- Level Set Methods

- Used to detect the border of an object on an image
- They use partial differential equations to evolve the curve in the image
- Positive = object
- Negative = not object

Dynamic Parallelism



- Level Set Algorithm Implementation



```
checkCuda(cudaMemcpy(gpu.intensity,
intensity,
gpu.size*sizeof(int),
cudaMemcpyHostToDevice));

checkCuda(cudaMemcpy(gpu.labels,
labels,
gpu.size*sizeof(int),
cudaMemcpyHostToDevice));

checkCuda(cudaDeviceSynchronize());

#ifdef(CUDA_TIMING)
float Ktime;
TIMER_CREATE(Ktime);
TIMER_START(Ktime);
#endif

#ifdef(VERBOSE)
printf("Running algorithm on GPU.\n");
#endif

// Launch kernel to begin image segmentation
evolveContour<<<1, numLabels>>>(gpu.intensity,
gpu.labels,
gpu.phi,
gpu.phiOut,
gridXSize,
gridYSize,
gpu.targetLabels,
gpu.lowerIntensityBounds,
gpu.upperIntensityBounds,
max_iterations,
gpu.globalBlockIndicator,
gpu.globalFinishedVariable,
gpu.totalIterations);


checkCuda(cudaDeviceSynchronize());

#ifdef(CUDA_TIMING)
TIMER_END(Ktime);
printf("Kernel Execution Time: %f ms\n", Ktime);
#endif
```


Dynamic Parallelism



- Level Set Algorithm Implementation



```
checkCuda(cudaMemcpy(gpu.intensity,
    intensity,
    gpu.size*sizeof(int),
    cudaMemcpyHostToDevice));
```

```
checkCuda(cudaMemcpy(gpu.labels,
    labels,
    gpu.size*sizeof(int),
    cudaMemcpyHostToDevice));
```

```
checkCuda(cudaDeviceSynchronize());
```

```
#if defined(CUDA_TIMING)
    float Ktime;
    TIMER_CREATE(Ktime);
    TIMER_START(Ktime);
#endif
```

```
#if defined(VERBOSE)
    printf("Running algorithm on GPU.\n");
#endif
```



```
// Launch kernel to begin image segmentation
evolveContour<<<1, numLabels>>>(gpu.intensity,
    gpu.labels,
    gpu.phi,
    gpu.phiOut,
    gridSize,
    gridYSize,
    gpu.targetLabels,
    gpu.lowerIntensityBounds,
    gpu.upperIntensityBounds,
    max_iterations,
    gpu.globalBlockIndicator,
    gpu.globalFinishedVariable,
    gpu.totalIterations);
```

```
#if defined(CUDA_TIMING)
    TIMER_END(Ktime);
    printf("Kernel Execution Time: %f ms\n", Ktime);
#endif
```

Dynamic Parallelism



- Level Set Algorithm Implementation



```
_global_ void evolveContour(unsigned int* intensity,
                           unsigned int* labels,
                           signed int* phi,
                           signed int* phiOut,
                           int gridSize,
                           int gridSize,
                           int* targetLabels,
                           int* lowerIntensityBounds,
                           int* upperIntensityBounds,
                           int max_iterations,
                           int* globalBlockIndicator,
                           int* globalFinishedVariable,
                           int* totalIterations ) {

    int tid = threadIdx.x;

    // Setting up streams for
    cudaStream_t stream;
    cudaStreamCreateWithFlags (&stream, cudaStreamNonBlocking);

    // Total iterations
    totalIterations = &totalIterations[tid];

    // Size in ints
    int size = (gridSize*BLOCK_TILE_SIZE)*(gridSize*BLOCK_TILE_SIZE);

    // New phi pointer for each label.
    phi = &phi[tid*size];
    phiOut = &phiOut[tid*size];

    globalBlockIndicator = &globalBlockIndicator[tid*gridSize*gridSize];

    // Global synchronization variable
    globalFinishedVariable = &globalFinishedVariable[tid];

    dim3 dimGrid(gridSize, gridSize);
    dim3 dimBlock(BLOCK_TILE_SIZE, BLOCK_TILE_SIZE);

    // Initialize phi array
    lssStep1<<<dimGrid, dimBlock, 0, stream>>>(intensity,
                                                labels,
                                                phi,
                                                targetLabels[tid],
                                                lowerIntensityBounds[tid],
                                                upperIntensityBounds[tid],
                                                globalBlockIndicator);

    int iterations = 0;
    do {
        iterations++;
        lssStep2<<<dimGrid, dimBlock, 0, stream>>>(phi,
                                                    globalBlockIndicator,
                                                    globalFinishedVariable );
        cudaDeviceSynchronize();
    } while (atomicExch(globalFinishedVariable,0) && (iterations < max_iterations));




    lssStep3<<<dimGrid, dimBlock, 0, stream>>>(phi,
                                                phiOut);

    *totalIterations = iterations;
}
```

Dynamic Parallelism



- Level Set Algorithm Implementation



```
_global_ void evolveContour(unsigned int* intensity,
                           unsigned int* labels,
                           signed int* phi,
                           signed int* phiOut,
                           int gridSize,
                           int gridSize,
                           int* targetLabels,
                           int* lowerIntensityBounds,
                           int* upperIntensityBounds,
                           int max_iterations,
                           int* globalBlockIndicator,
                           int* globalFinishedVariable,
                           int* totalIterations ) {

    int tid = threadIdx.x;

    // Setting up streams for
    cudaStream_t stream;
    cudaStreamCreateWithFlags (&stream, cudaStreamNonBlocking);

    // Total iterations
    totalIterations = &totalIterations[tid];

    // Size in ints
    int size = (gridSize*BLOCK_TILE_SIZE)*(gridSize*BLOCK_TILE_SIZE);

    // New phi pointer for each label.
    phi = &phi[tid*size];
    phiOut = &phiOut[tid*size];

    globalBlockIndicator = &globalBlockIndicator[tid*gridSize*gridSize];

    // Global synchronization variable
    globalFinishedVariable = &globalFinishedVariable[tid];

    dim3 dimGrid(gridSize, gridSize);
    dim3 dimBlock(BLOCK_TILE_SIZE, BLOCK_TILE_SIZE);

    // Initialize phi array
    lssStep1<<<dimGrid, dimBlock, 0, stream>>>(intensity,
                                                labels,
                                                phi,
                                                targetLabels[tid],
                                                lowerIntensityBounds[tid],
                                                upperIntensityBounds[tid],
                                                globalBlockIndicator);

    int iterations = 0;
    do {
        iterations++;
        lssStep2<<<dimGrid, dimBlock, 0, stream>>>(phi,
                                                    globalBlockIndicator,
                                                    globalFinishedVariable );
        cudaDeviceSynchronize();
    } while (atomicExch(globalFinishedVariable,0) && (iterations < max_iterations));
    lssStep3<<<dimGrid, dimBlock, 0, stream>>>(phi,
                                                phiOut);

    *totalIterations = iterations;
}
```

- To understand recursive usage
 - Classic Fibonacci series: 0, 1, 1, 2, 3, 5...
 - CPU Recursive
 - GPU non recursive nonDP
 - GPU recursive DP

Fibonacci – CPU Recursion



```
int fib(int n){  
    if (n == 0 || n == 1){  
        return n;  
    }  
    else{  
        return fib(n-1) + fib(n-2);  
    }  
}
```

Fibonacci – GPU Basic Kernel




```
__global__ void fib_kernel_plain(int n, long int* vFib){  
    int tid = threadIdx.x + blockDim.x * blockIdx.x;  
    if (tid > n/32)  
        return;  
  
    if (n == 0 || n == 1){  
        return;  
    }  
    for(int i=tid*32 + 2; i <= n && i < tid*32 + 32; i++){  
        vFib[i] = vFib[i-1] + vFib[i-2];  
    }  
}
```

Fibonacci – GPU Recursion Kernel

```
__global__ void fib_kernel_par_rec(int n, unsigned long int* vFib){  
    if (n == 0 || n == 1)  
        return;  
  
    fib_kernel_par_rec<<<1, 1>>>(n-2, vFib);  
    fib_kernel_par_rec<<<1, 1>>>(n-1, vFib);  
    cudaDeviceSynchronize();  
    vFib[n] = vFib[n-1] + vFib[n-2];  
  
}
```


Results

- Performance CPU vs GPU non DP vs GPU DP




Results

- Performance CPU vs GPU non DP vs GPU DP



Results

- Performance CPU vs GPU non DP vs GPU DP



What is and what is not dynamic parallelism



- CDP ensures better work balance, and offers advantages in terms of programmability
 - However, launching grids with a very small number of threads could lead to severe underutilization of the GPU resources
 - A general recommendation
 - Child grids with a large number of thread blocks,
 - Or at least thread blocks with hundreds of threads, if the number of blocks is small
 - Nested parallelism for tree processing
 - Thick tree nodes (each node deploys many threads) work well
 - And/or when branch degree is large (each parent node has many children)
 - As the nesting depth is limited in hardware, only relatively shallow trees can be implemented efficiently.
-